

# Lec 1 Hello GPU

Tonghua Su  
School of Software  
Harbin Institute of Technology

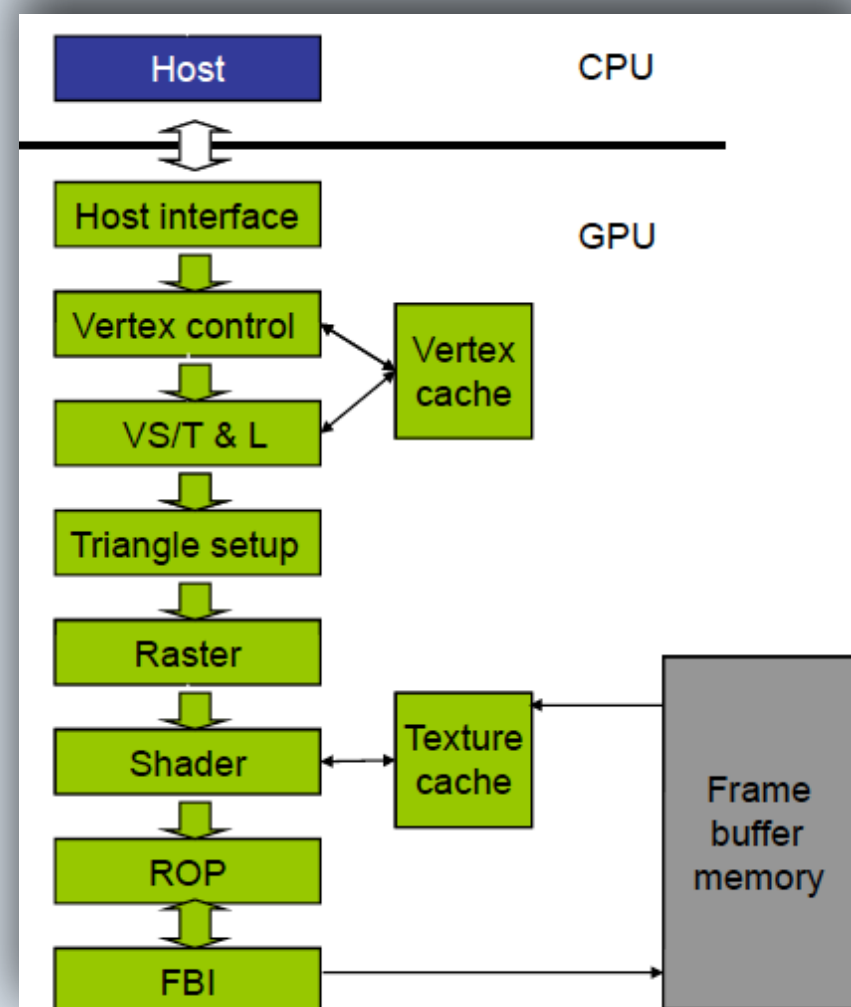

# Outline

- 1 **What is GPU?**
- 2 **GPU Architecture (briefly)**
- 3 **First GPU Program**
- 4 **Amdahl's Law**
- 5 **HW1**

# What is GPU

- Specialized electronic circuit


- ✓ Accelerate the building of images intended for display



```
0.748952 -0.764952 -0.210132,  
0.872245 -0.600062 -0.210132,  
1.00016 -0.369596 -0.210132,  
1.00939 -0.004004 -0.210132,  
1.14496 0.324436 -0.210132,  
1.15747 0.601712 -0.210132,  
1.08016 0.793529 -0.210132,  
0.09164 0.972032 -0.210132,  
0.808203 0.929016 -0.210132,  
0.442563 0.985585 -0.210132,  
0.221794 1.00159 -0.210132,  
0 1.0053 -0.210132,  
-0.221794 1.00159 -0.210132,  
-0.442563 0.985585 -0.210132,  
-0.808203 0.929016 -0.210132,
```

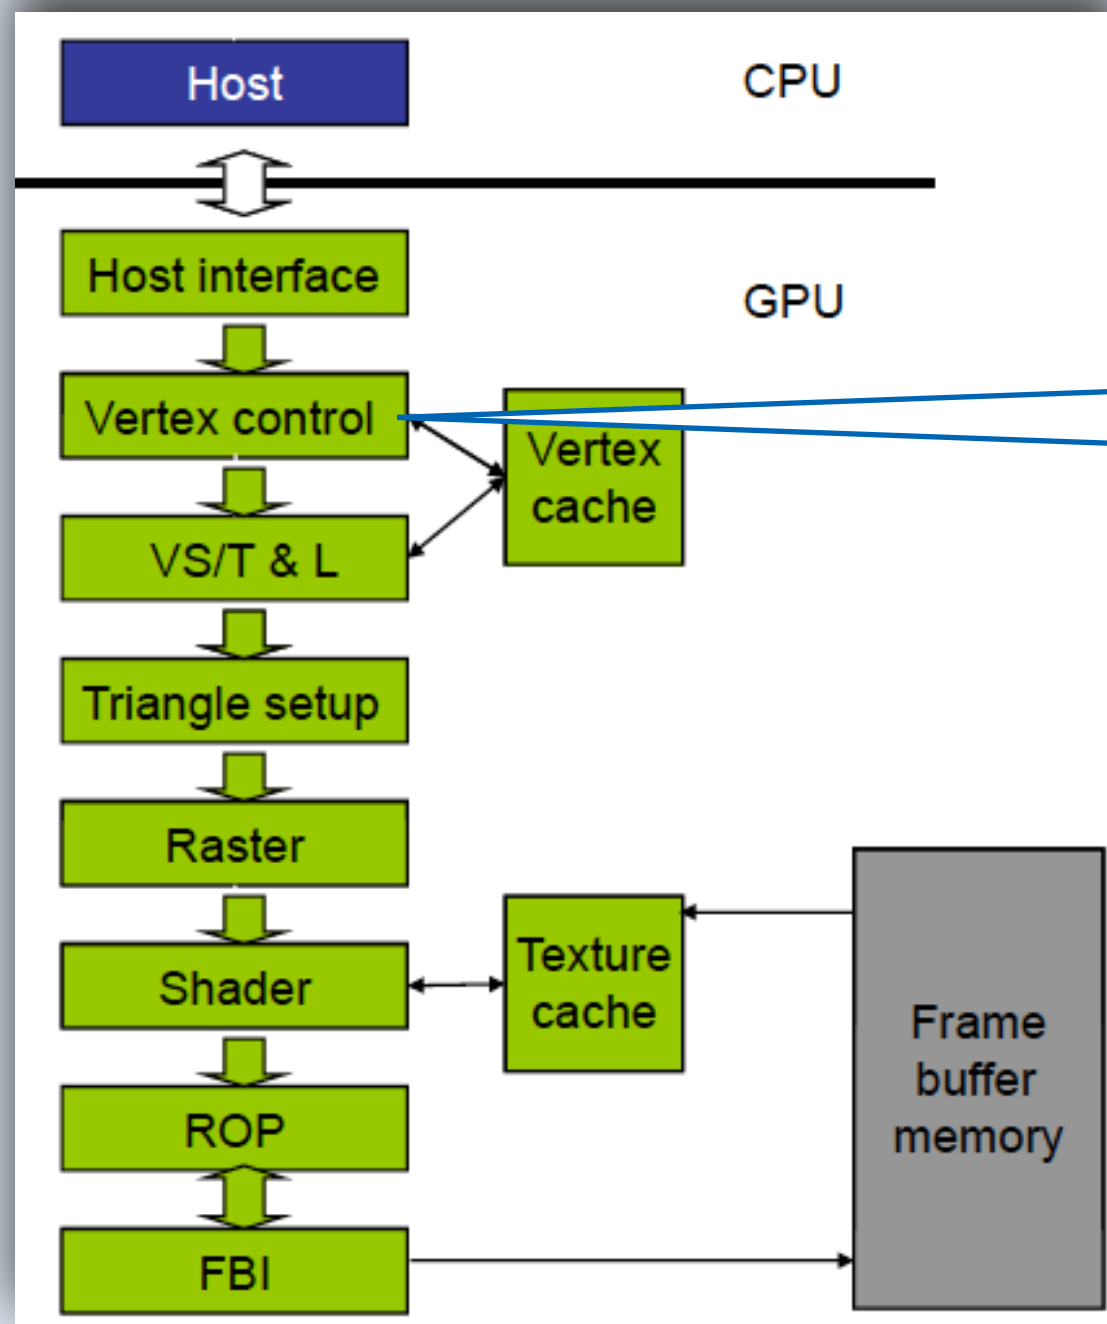


# What is GPU



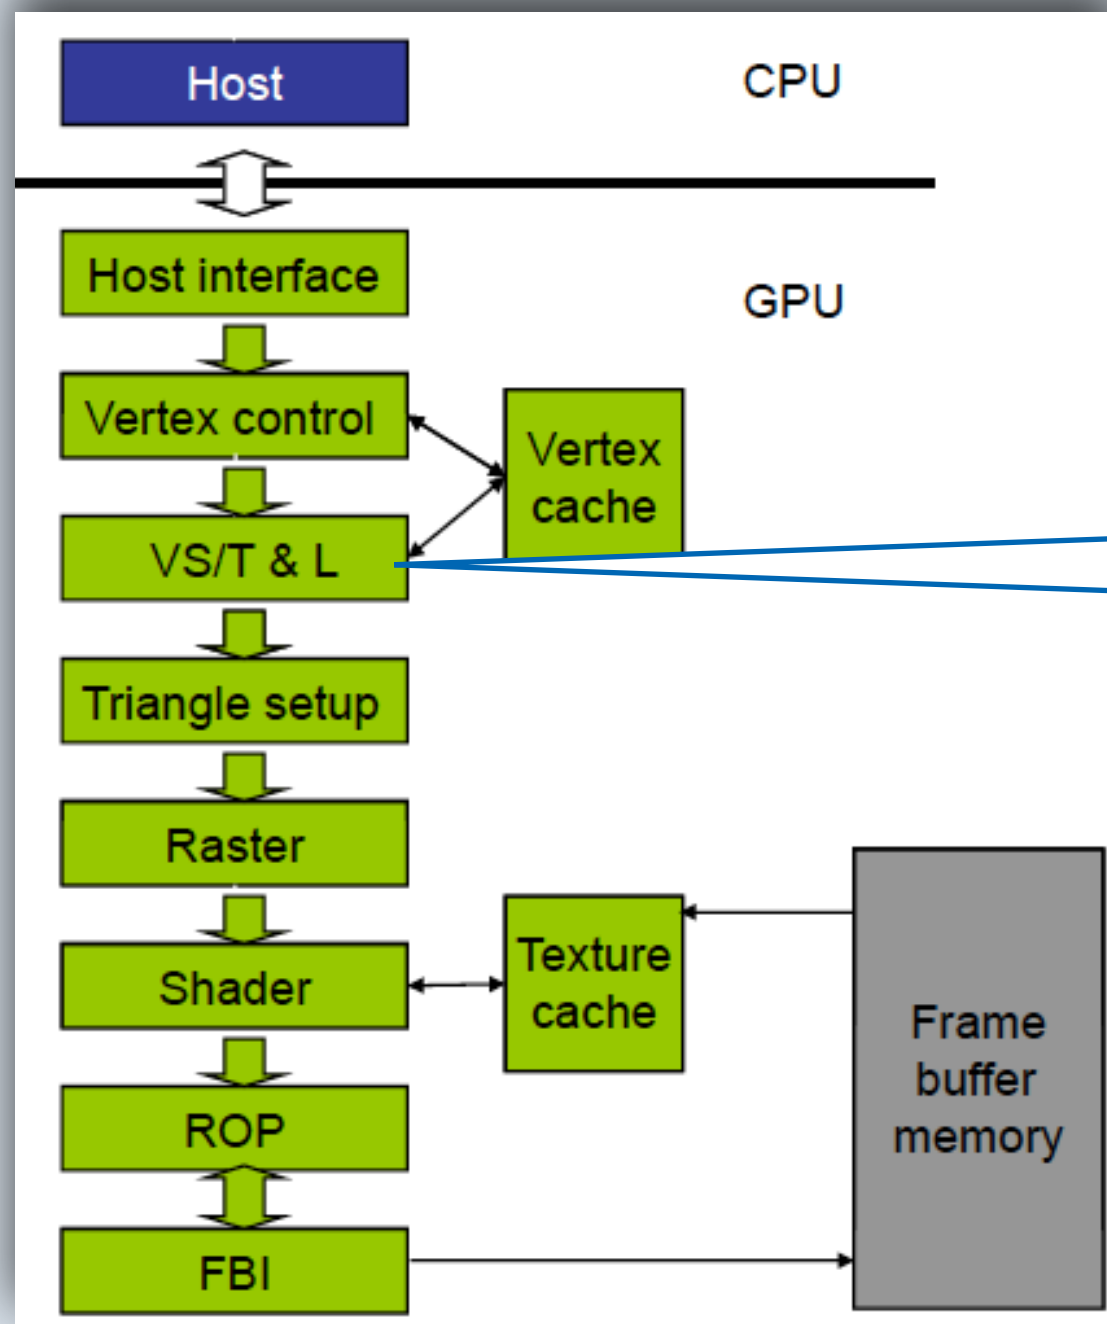
Receives graphics commands and data from CPU

# What is GPU



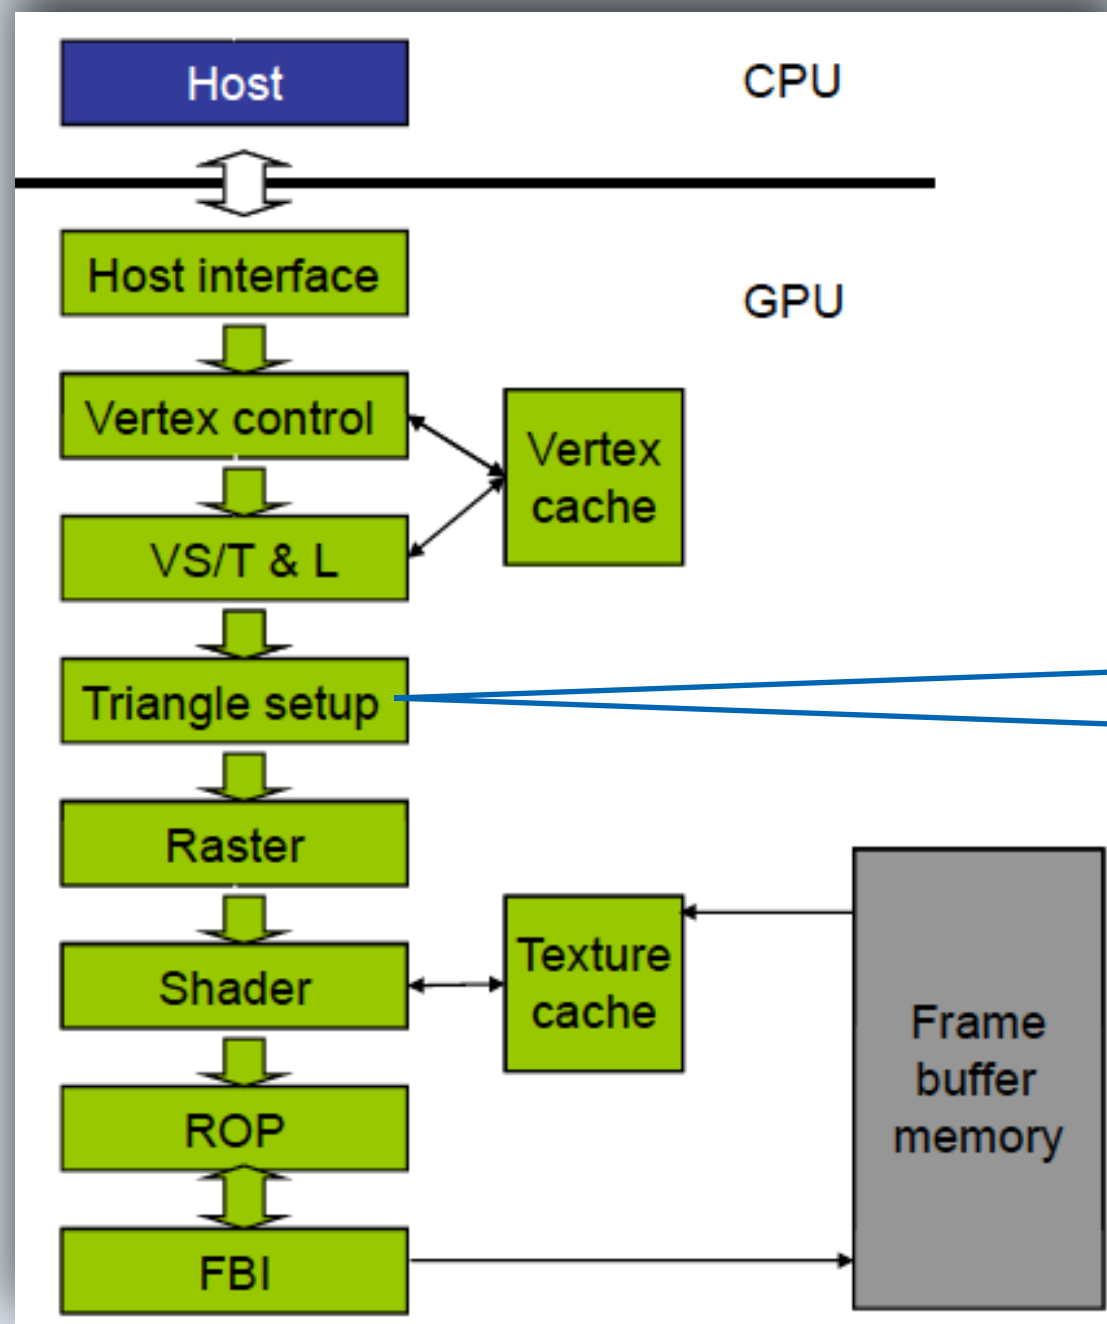
Receives triangle data  
Converts them into a form that hardware understands  
Store the prepared data in vertex cache

# What is GPU



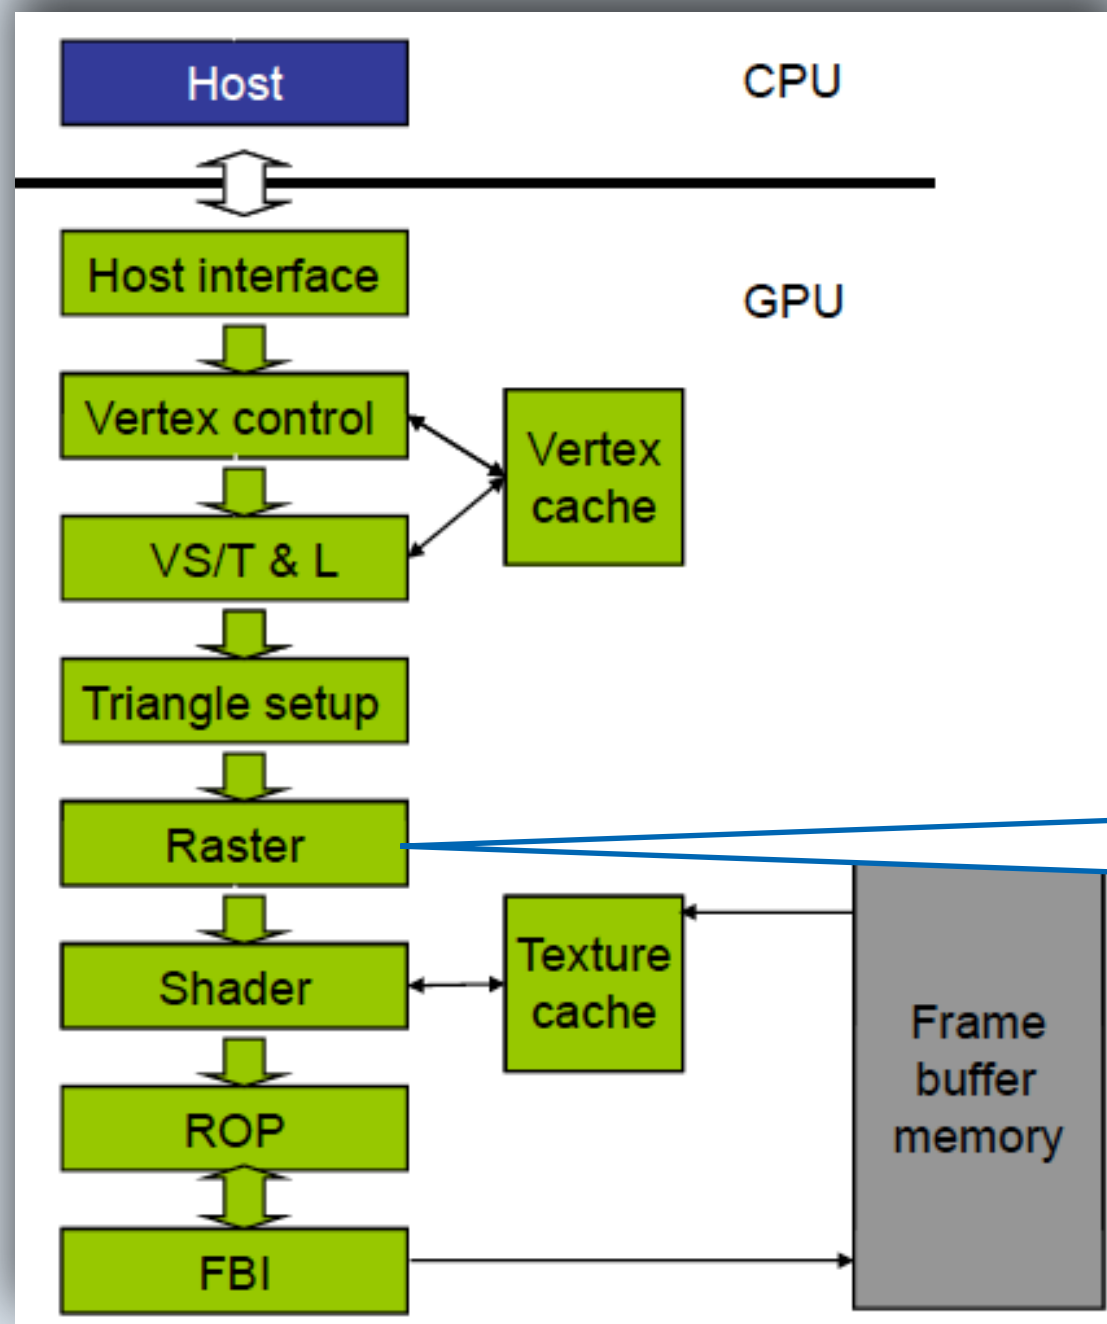
- Vertex shading transform and lighting
- Assigns per-vertex value

# What is GPU



Creates edge equations to interpolate colors across pixels touched by the triangle

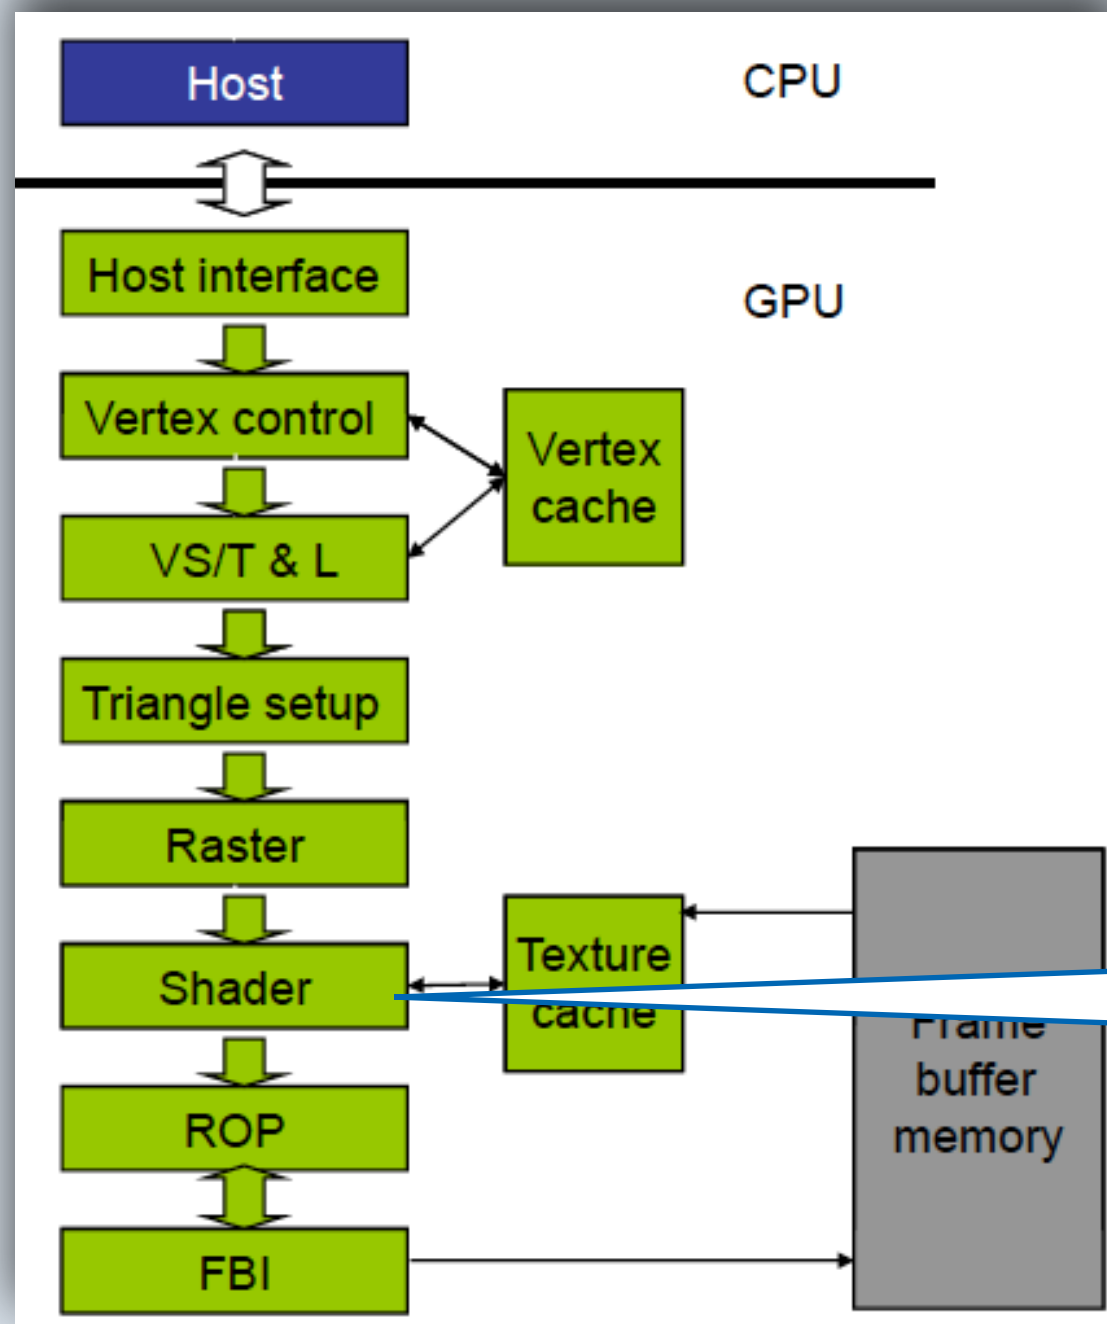
# What is GPU



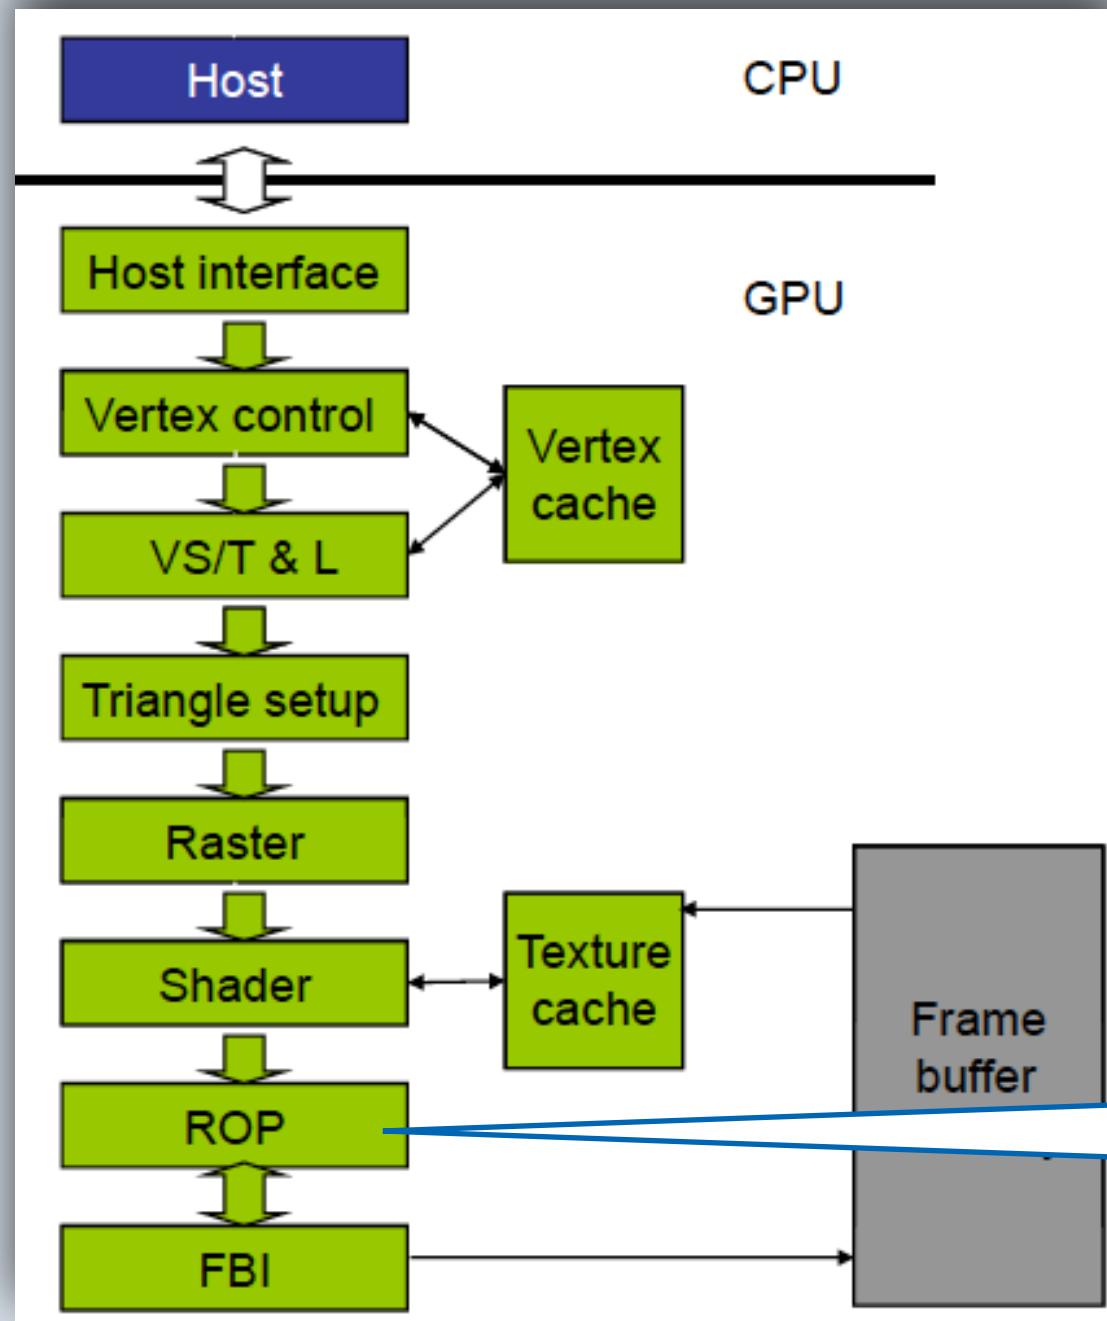
- Determines which pixel falls in which triangle
- For each pixel, interpolate per-pixel values



# What is GPU

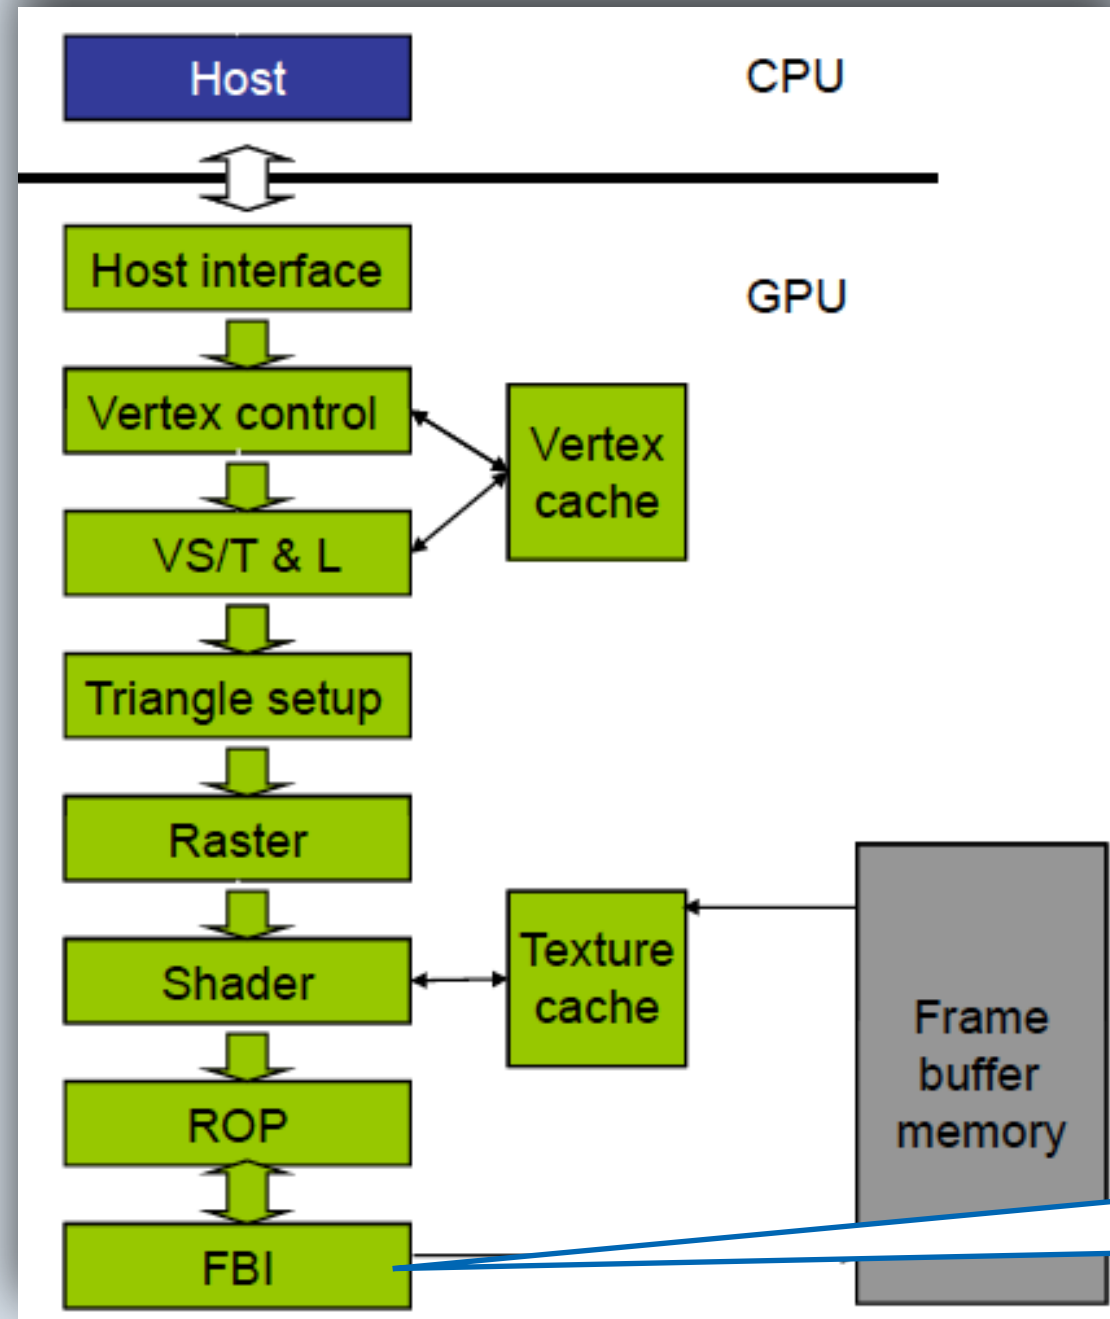


# What is GPU



Determine The raster operation: performs color raster operations that blend the color of overlapping objects for transparency and antialiasings the final color of each pixel

# What is GPU




The frame buffer interface manages memory reads/writes


# What is GPU

- **Specialized electronic circuit**


- ✓ Accelerate the building of images intended for display



Triangle Geometry



Aliased



Anti-Aliased



- ✓ But now, it is used for general-purpose computation!

“Amper”  
?? xtors

# What is GPU

- High throughput computation

- ✓ GeForce GTX 980: 4,612 GFLOP/s

- High bandwidth memory

- ✓ GeForce GTX 980: 224 GB/s

- High availability to all

- ✓ 500+ million CUDA-capable GPUs in the wild

“Turing”  
?? xtors

“Volta”  
?? xtors

“Pascal”  
?? xtors

“Maxwell”  
5.2B xtors

“Kepler”  
3.5~7B xtors

“Fermi”  
3B xtors

GeForce 8800  
681M xtors

GeForce FX  
125M xtors

GeForce 3  
60M xtors

GeForce® 256  
23M xtors

RIVA 128  
3M xtors



# Flynn Taxonomy

## ●Flynn分类法(行为特征)

<b><i>S I S D</i></b> <b><i>Single Instruction, Single Data</i></b>  串行计算机( <b><i>von Neumann</i></b> 计算机)	<b><i>S I M D</i></b> <b><i>Single Instruction, Multiple Data</i></b>  特定领域的加速器( <b><i>GPU</i></b> 、向量处理机等)
<b><i>M I S D</i></b> <b><i>Multiple Instruction, Single Data</i></b>  比较少见，多用于容错计算机	<b><i>M I M D</i></b> <b><i>Multiple Instruction, Multiple Data</i></b>  常见的并行计算机都可归入此类 <b><i>MPP/Cluster/SMP/</i></b> 当前基于 <b><i>Cache</i></b> 的 <b><i>Multi-core (Intel、AMD)</i></b>

# Quiz


- **Pascal 架构、Volta架构与Turing架构有多少晶体管?**
- **对比CPU与GPU浮点计算能力**
  - **CPU以E5 2640 V3为例, ~300GFLOPS**
  - **GPU以GTX 980为例**

# Outline

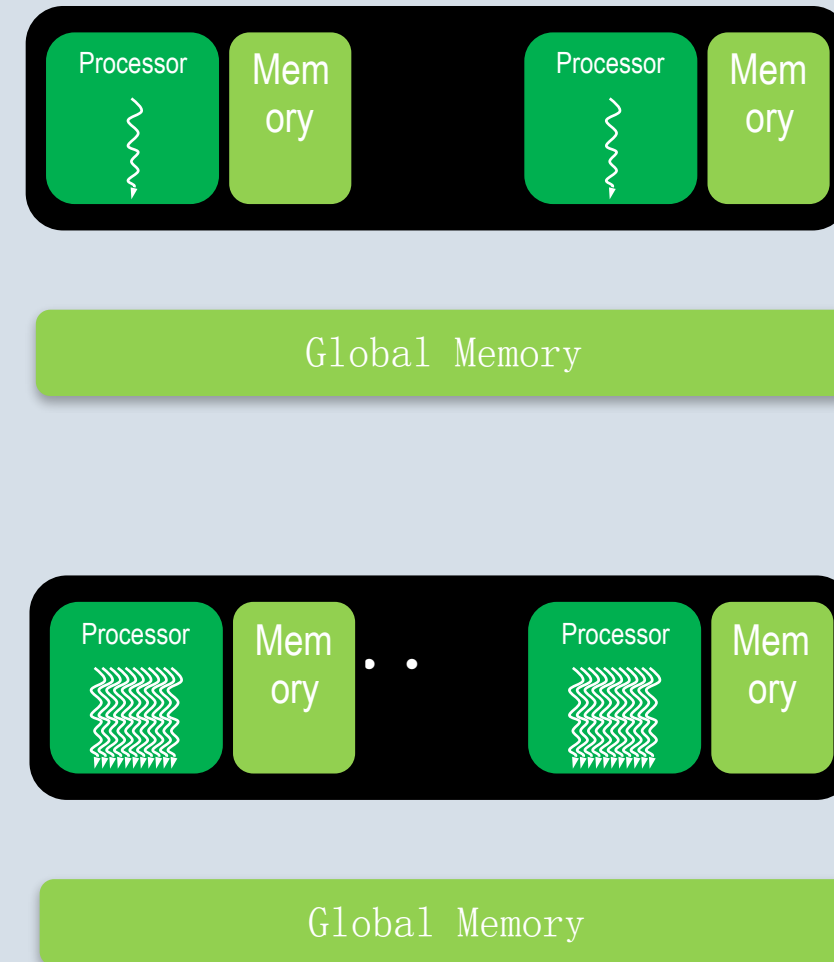
- 1 What is GPU?
- 2 GPU Architecture (briefly)
- 3 First GPU Program
- 4 Amdahl's Law
- 5 HW 1



# NVIDIA GPU Architecture



Fermi GF100




Multi-core

Many-core

# Streaming Multi-processor (SM)

- **32 CUDA Cores per SM (512 total)**
- **Direct load/store to memory**
  - ✓ Usual linear sequence of bytes
  - ✓ High bandwidth (Hundreds GB/sec)
- **64KB of fast, on-chip RAM**
  - ✓ Software or hardware-managed
  - ✓ Shared amongst CUDA cores
  - ✓ Enables thread communication



# HW 1.1 Device Query

## ● 查询你机器上GPU设备的参数


- ✓ 新建**.cu**文件
- ✓ 调用**cudaGetDeviceCount()**得到GPU设备的数量
- ✓ 调用**cudaGetDeviceProperties()**函数得到GPU设备的属性结构体
- ✓ 解释关键属性的含义，至少包括设备名称、计算能力为多少、设备可用全局内存、每线程块最大线程数、设备可用全局内存容量、每线程块可用共享内存容量、每线程块可用寄存器数量、每线程块最大线程数、每个处理器簇最大驻留线程数、设备中的处理器簇数量等
- ✓ 可参考**WILT 3.2**节

# Outline

- 1 **What is GPU?**
- 2 **GPU Architecture (briefly)**
- 3 **First GPU Program**
- 4 **Amdahl's Law**
- 5 **HW 1**

# Hello CUDA

## ●Vector Sum



## ●Serial code

```
for (i=0;i<128;i++)  
{  
    c[i] = a[i] + b[i];  
}
```

## ●Translate into CUDA threads


```
__global__ void addKernel(int * const a, const int * const b, const int * const c)  
{  
    c[i] = a[i] + b[i];  
}
```

# Hello CUDA

## ● Vector Sum

### ✓ Identify thread id

```
__global__ void addKernel(int * const a, const int * const b, const int * const c)
{
    const unsigned int i = threadIdx.x;
    c[i] = a[i] + b[i];
}
```



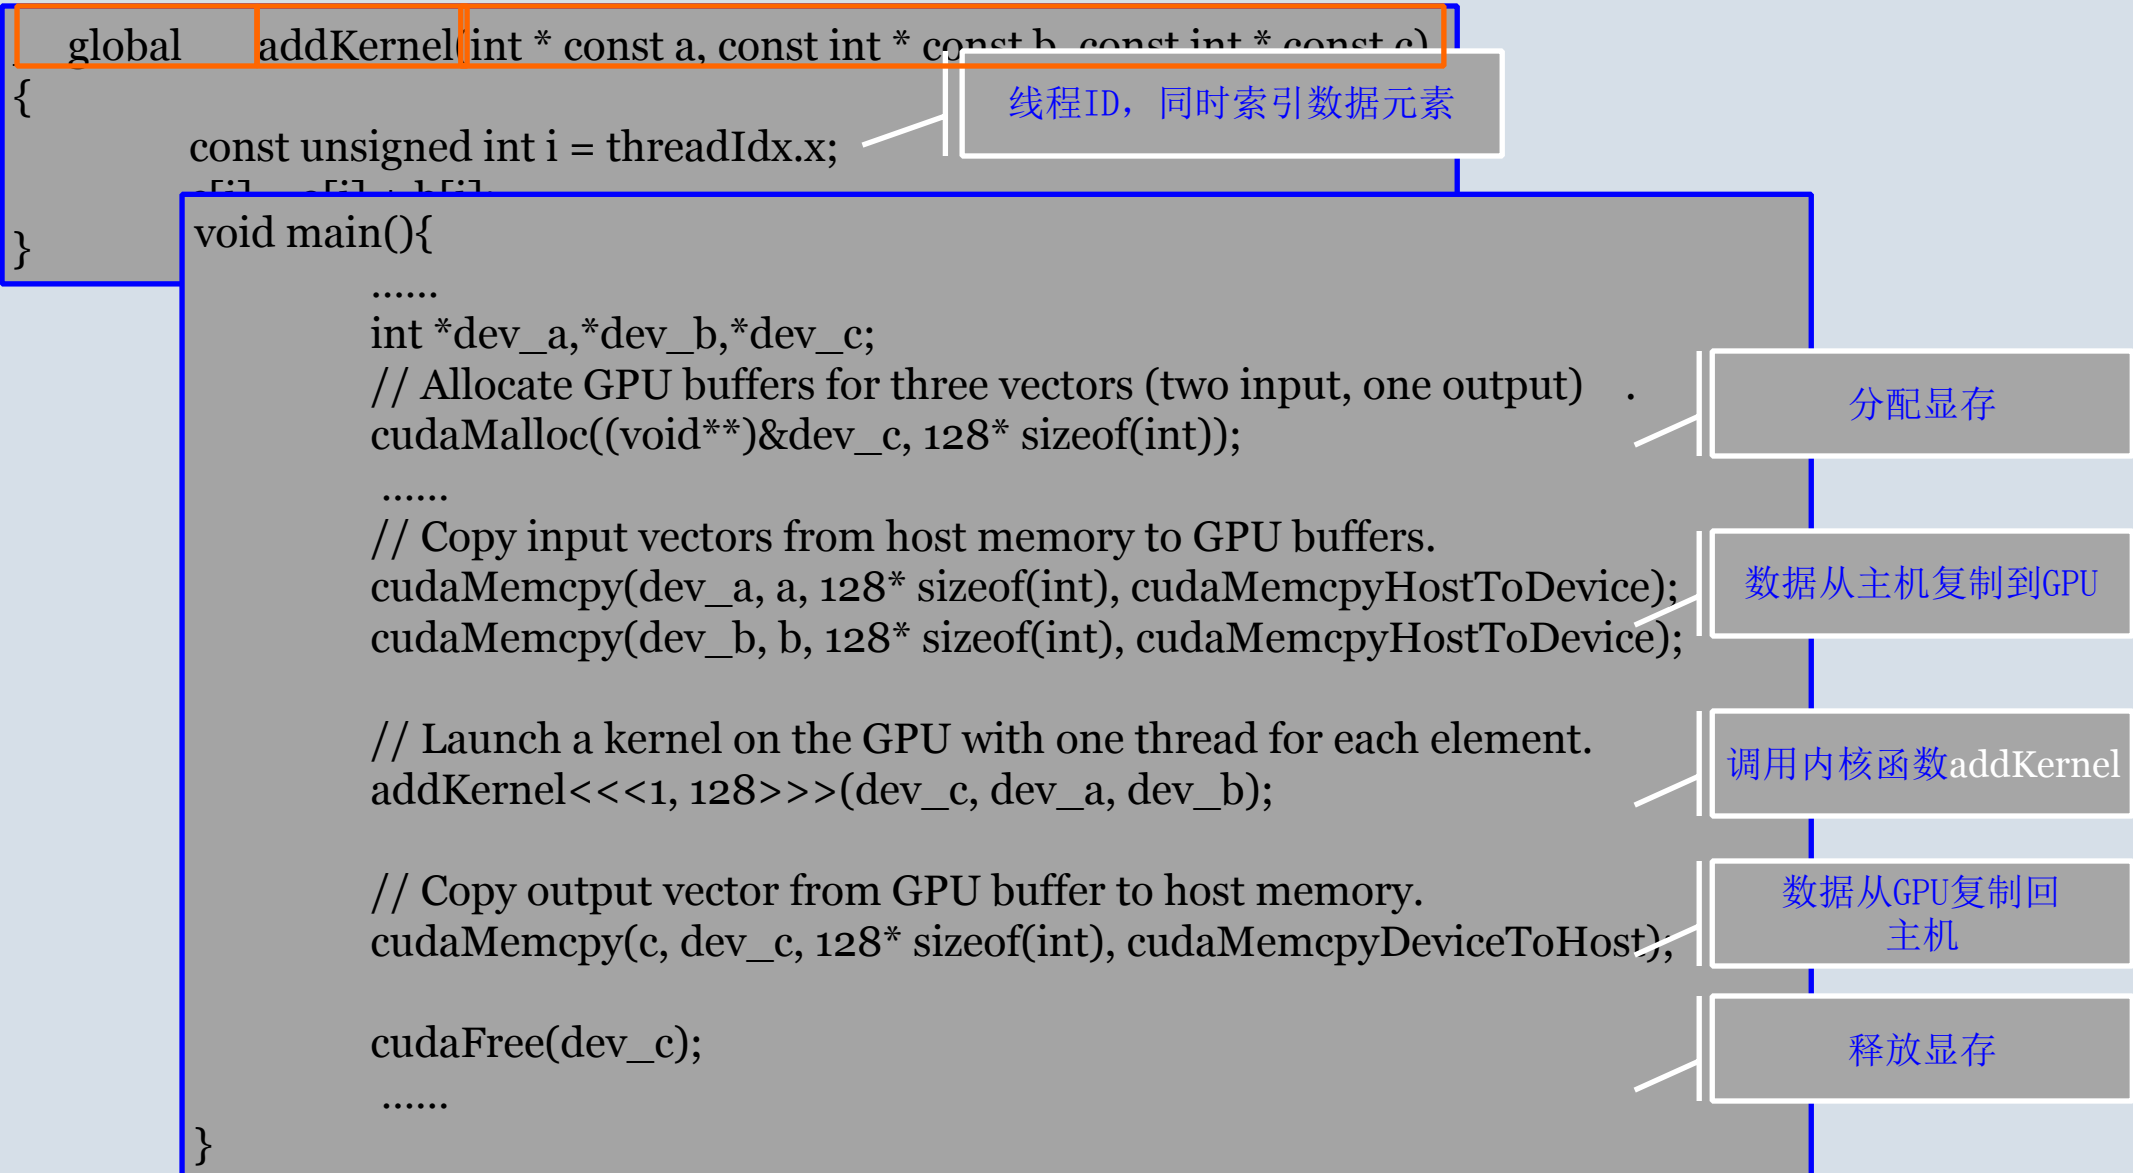
### ✓ Invoke CUDA kernel

- `kernel_function<<<num_blocks, num_threads>>>(param1, param2, .)`
- E.g. `addKernel<<< 1, 128 >>>(a, b, c);`

# Hello GPU

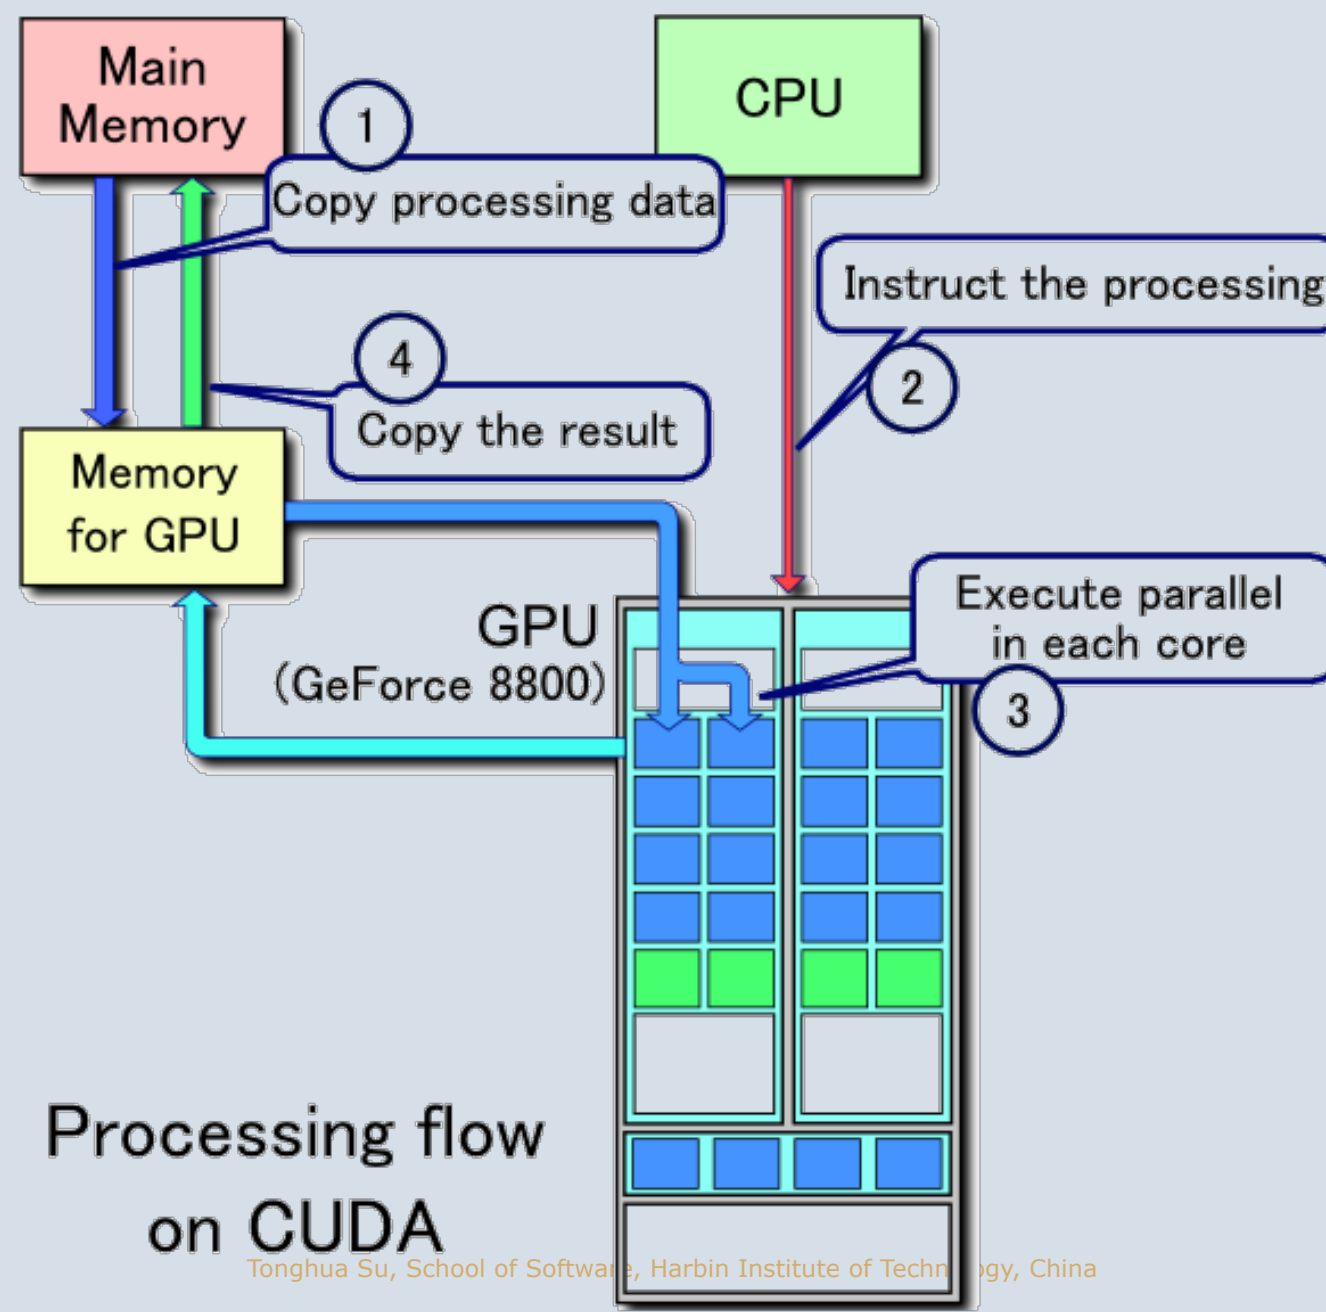
## ● Demo 1

# Hello CUDA





# Hello CUDA



Processing flow  
on CUDA

## HW 1.2 第一个GPU程序

- 自己从头编写并运行VectorSum内核
-

# Outline

- 1 What is GPU?
- 2 GPU Architecture (briefly)
- 3 First GPU Program
- 4 Amdahl's Law
- 5 HW 1

# Amdahl's Law

- Gene Amdahl in 1967:

$$Speedup = \frac{1}{r_s + \frac{r_p}{N}}$$

where  $r_s + r_p = 1$  and  $r_s$  represents the ratio of the sequential portion.

✓ Consider: 30% portion of time with 100X speedup through paralleling vs 99% portion with 100X

Amdahl, Gene (1967). "Validity of the Single Processor Approach to Achieving Large-Scale Computing Capabilities". *AFIPS Conference Proceedings* (30): 483–485.

# Quiz

● 请使用**Amdahl** 的公式计算下面程序的加速比

✓ 程序**99%**的部分可以使用无限多个核心计算，剩余的**1%**无法并行

# Outline

- 1 What is GPU?
- 2 GPU Architecture (briefly)
- 3 First GPU Program
- 4 Amdahl's Law
- 5 HW 1

# HW 1

## ●1.1 查询你机器上GPU设备的参数

- ✓ 新建`.cu`文件
- ✓ 调用`cudaGetDeviceCount()`得到GPU设备的数量
- ✓ 调用`cudaGetDeviceProperties()`函数得到GPU设备的属性结构体
- ✓ 解释关键属性的含义，至少包括设备名称、计算能力为多少、设备可用全局内存、每线程块最大线程数、设备可用全局内存容量、每线程块可用共享内存容量、每线程块可用寄存器数量、每线程块最大线程数、每个处理器簇最大驻留线程数、设备中的处理器簇数量等
- ✓ 可参考WILT 3.2节

## ●1.2 自己从头编写并运行VectorSum内核