

# Review on Lec 1 & Lec 2

## Hello GPU

```
__global__ void addKernel(int * const a, const int * const b, const int *
const c)
{
    const unsigned int i = threadIdx.x;
    c[i] = a[i] + b[i];
}
```

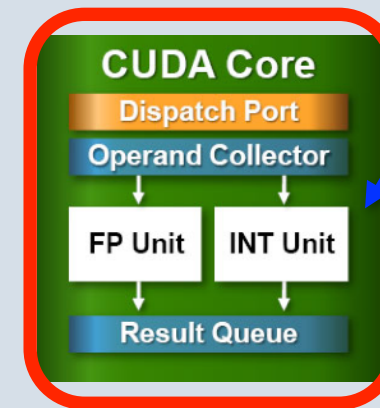
```
void main(){
    .....
    int *dev_a,*dev_b,*dev_c;
    cudaMalloc((void**)&dev_c, 128* sizeof(int));
    .....
    cudaMemcpy(dev_a, a, 128* sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b, 128* sizeof(int), cudaMemcpyHostToDevice);

    // Launch a kernel on the GPU with one thread for each element.
    addKernel<<<1, 128>>>(dev_c, dev_a, dev_b);

    cudaMemcpy(c, dev_c, 128* sizeof(int), cudaMemcpyDeviceToHost);

    cudaFree(dev_c);
    .....
}
```

## Kepler SMX



# **Lec 3** CUDA Software Abstraction

**Tonghua Su**  
School of Software  
Harbin Institute of Technology

# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# Outline


- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# Design Philosophy

**CPU: Latency Oriented Cores**

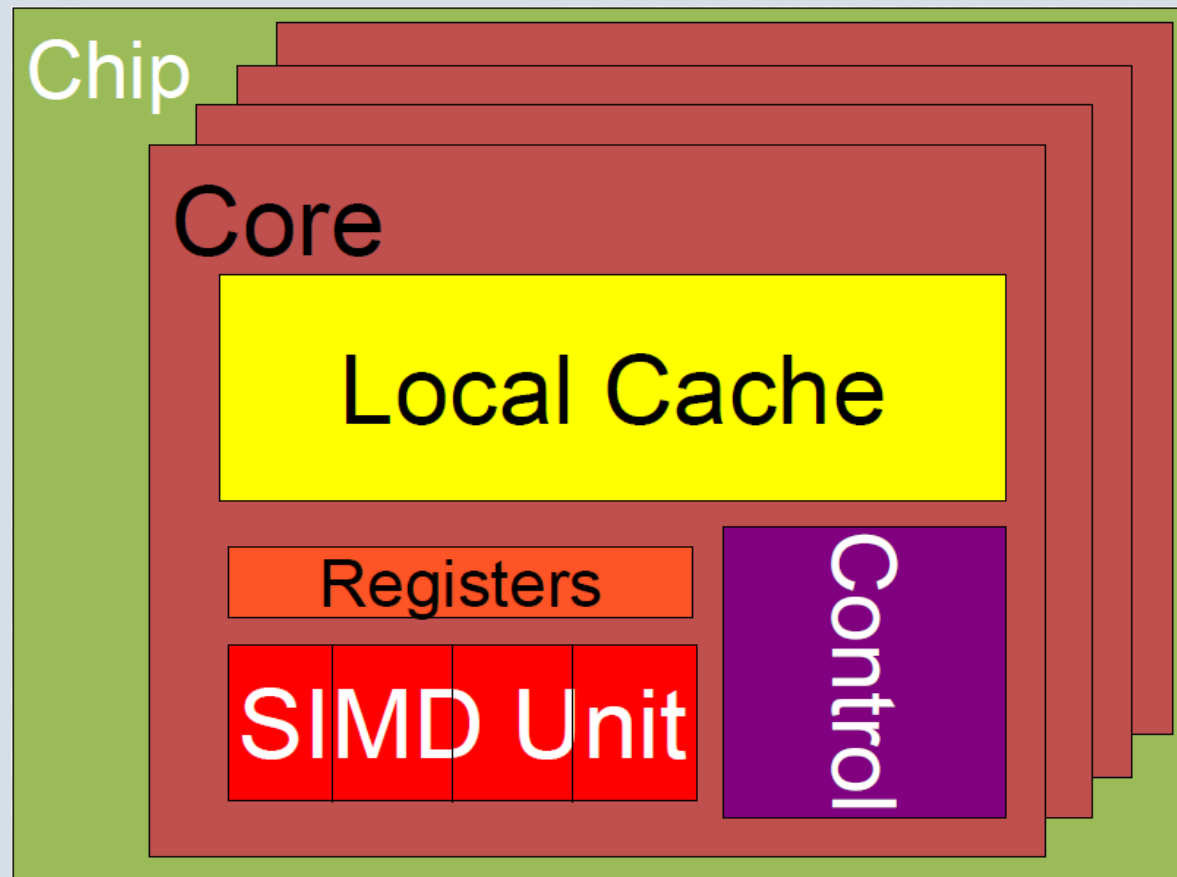


**GPU: Throughput Oriented Cores**

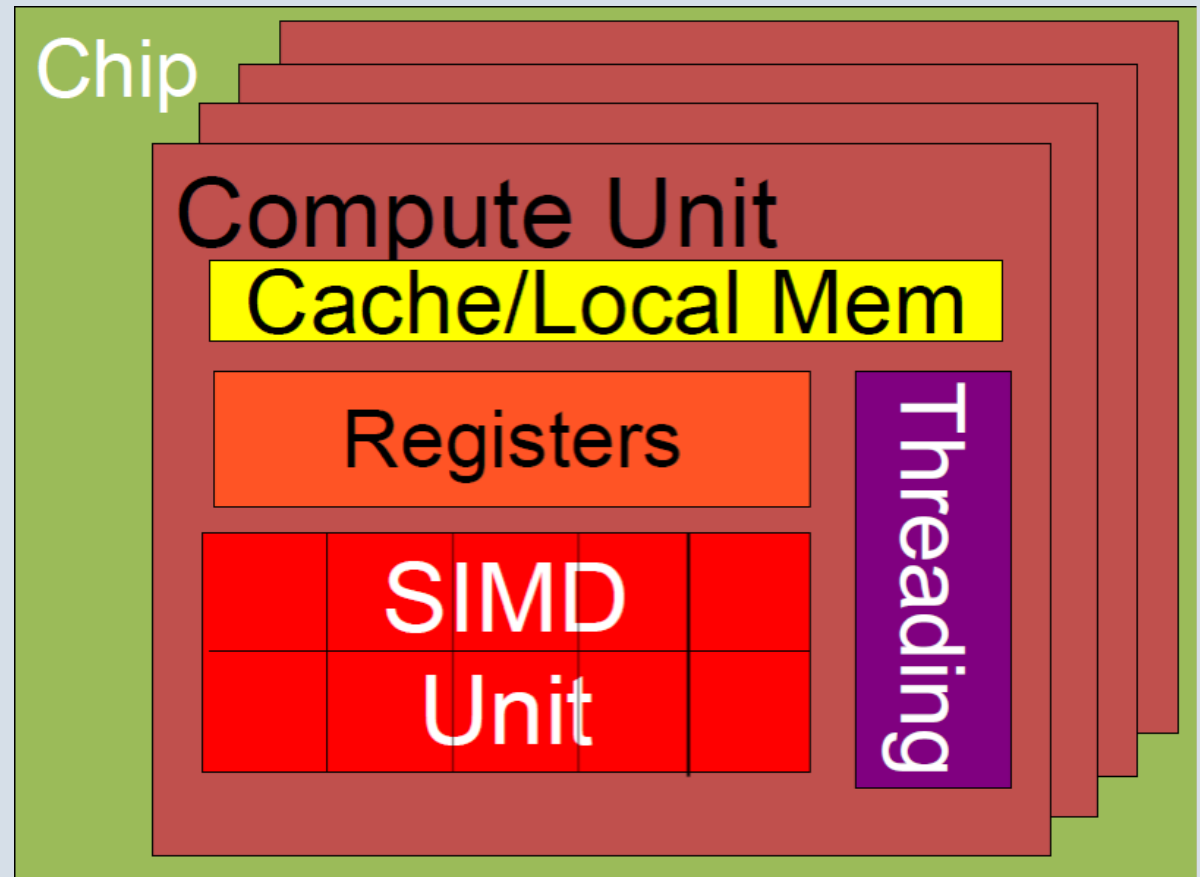


# Design Philosophy

## CPU: Latency Oriented Cores



## GPU: Throughput Oriented Cores



# Multithreading

## ●SIMD in CPU

- ✓ All cores execute the same instructions simultaneously, but with different data
- ✓ Similar to vector computing on CRAY supercomputers
- ✓ e. g. SSE4, AVX instruction set

## ●SIMT in SMX

- ✓ Multithreaded CUDA core
- ✓ Threads on each SMX execute in group sharing same instruction
- ✓ Fine-grained parallelism
- ✓ Natural for graphics processing and much scientific computing
- ✓ SIMT is also a natural choice for many-core chips to simplify each core

# Multithreading

- **Thread: instruction stream with own PC and data**
  - ✓ **Owning private register, private memory, program counter and thread execution state**
  - ✓ **Thread Level Parallelism(TLP): Exploit the parallelism inherent between threads**
- **Multithreading**
  - ✓ **Multiple threads to share the functional units of 1 processor via overlapping**
  - ✓ **Processor must duplicate independent state of each thread**
    - e.g., a separate copy of register file, a separate PC
  - ✓ **Often, hardware for fast thread switch**
  - ✓ **Memory to be shared**
  - ✓ **Solving the memory access stall**



# GPU Multithreading

## ● Multithreaded Hardware

- ✓ **CUDA core is multithreaded processor**

- Supporting 96 threads in GTX 8800

- ✓ **Run in group of 32 threads (called a warp)**

- ✓ **Zero-cost “context switching”**

- Each thread has its own registers (which limits the number of active threads)

- ✓ **Shared memory/L1 cache**

## ● Fine-grained multithreading

- ✓ **Able to switch between warps on each instruction**

- ✓ **Schedule without pre-emption**

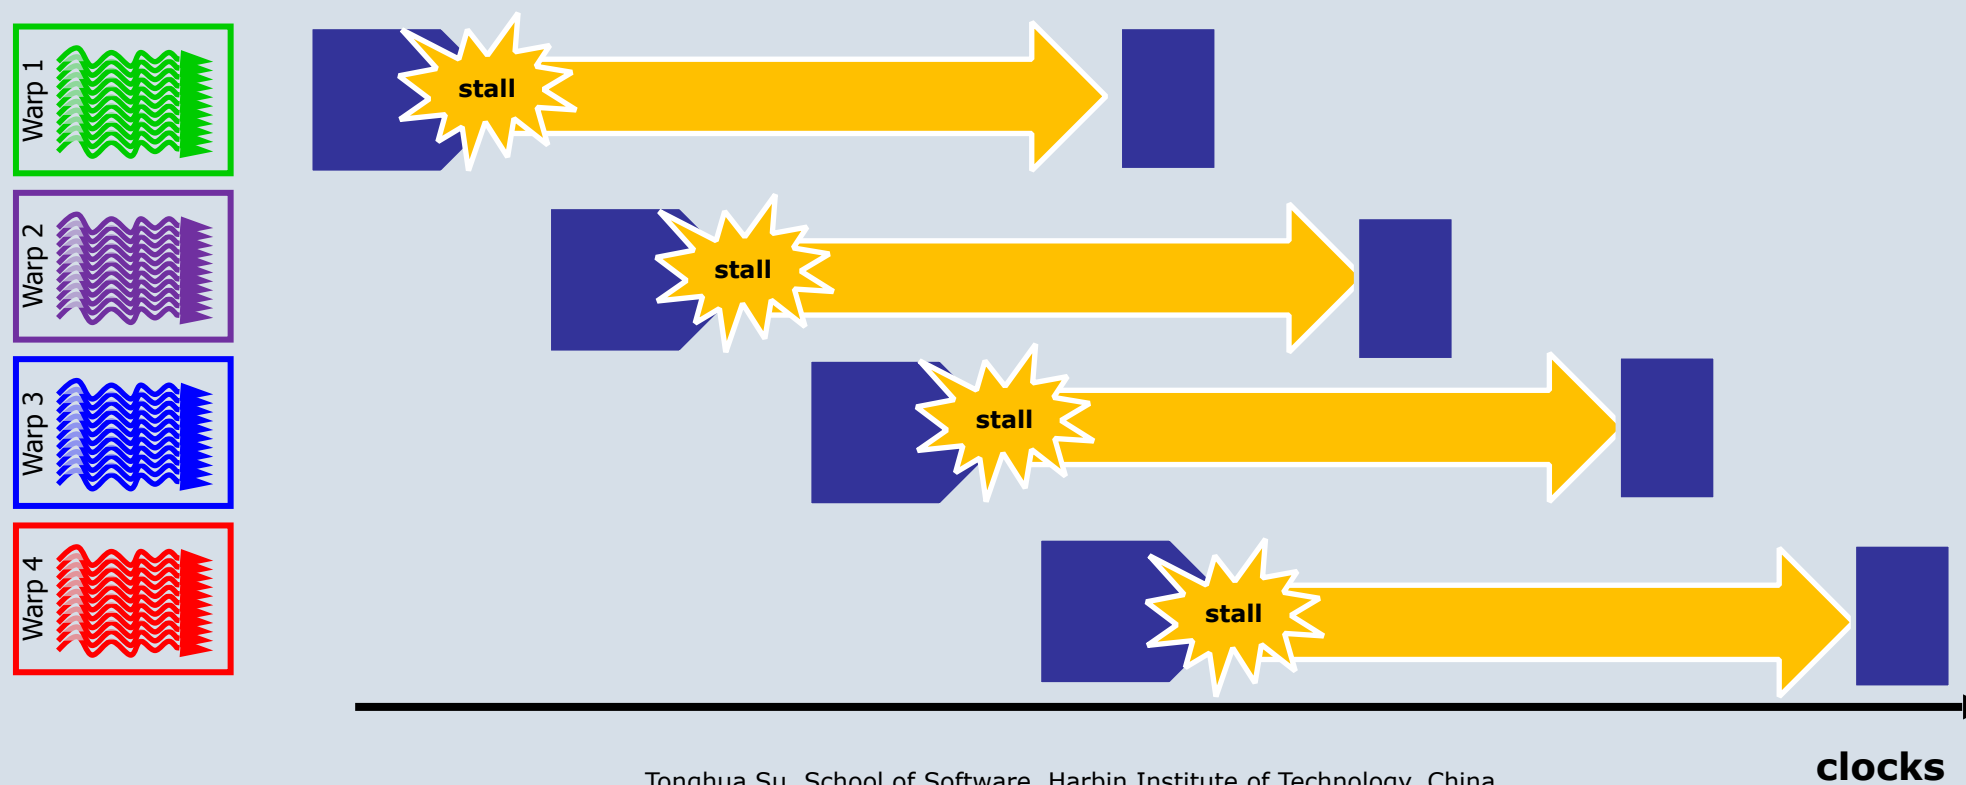
- ✓ **Scheduling eligible warps in a round-robin fashion**



# GPU Multithreading

## ●Hiding Latency Stalls

- ✓ Fetch → Decode → Execute → **Memory** → Writeback
- ✓ Execution alternates between “active” warps, with warps becoming temporarily “inactive” when waiting for data
- ✓ Lots of active warps is the key to high performance



# Quiz

- 定量计算内存访问带来的停滞
  - ✓ 假设从显存读取数据的延时是 **400**时钟周期
  - ✓ 如果每个线程束 (**warp**) 可以运行**10**个周期，那么另外需要多少个就绪的线程束才能掩盖停滞带来的时间缝隙?

# 作业预热

- 请编写程序，实现两个矩阵相乘
  - 先编写**CPU**版程序，然后给出**GPU**代码
  - 每个矩阵用线性数组表示
  - 考虑多个**block**
  - 考虑矩阵尺寸不是**block**尺寸的整倍数

# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# CUDA

- **CUDA(Compute Unified Device Architecture) is developed by Nvidia around 2007**
  - ✓ 2-4 week learning curve for those with experience of OpenMP and MPI programming
  - ✓ large user community on NVIDIA forums
- **CUDA is a parallel computing platform and programming model that enables dramatic increases in computing performance by harnessing the power of the GPU**

# CUDA

## ●CUDA as Parallel Computing Platform

- ✓ **Language:** CUDA C which based on C with some extensions(extensive C++ support)
- ✓ **Editor:** Eclipse/Visual Studio
- ✓ **Compiler:** nvcc
- ✓ **SDK:** CUDA toolkit, Libraries, Samples
- ✓ **Profiler & Debugger:** Nsight

## ●CUDA as Programming Model

- ✓ **Software Abstraction of GPU hardwares**
- ✓ **Independent to OSs, CPUs, Nvidia GPUs**


# CUDA Abstraction

## ●CUDA Virtualizes the Physical Hardware

- ✓ thread is a virtualized CUDA cores (registers, PC, state)
- ✓ block is a virtualized streaming multiprocessor (threads, shared mem.)

## ●Scheduled onto Physical Hardware without Pre-emption

- ✓ threads/blocks launch & run to completion/suspension
- ✓ blocks should be independent

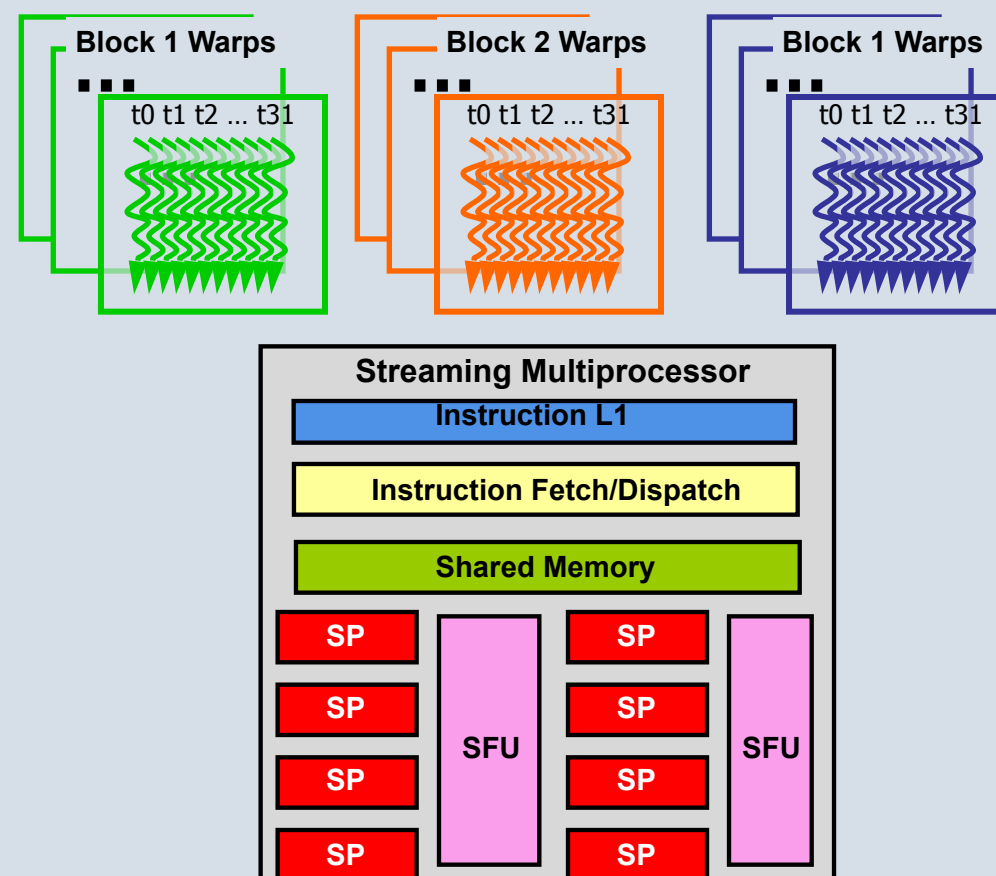




# CUDA Abstraction

## •Key Parallel Abstractions in CUDA

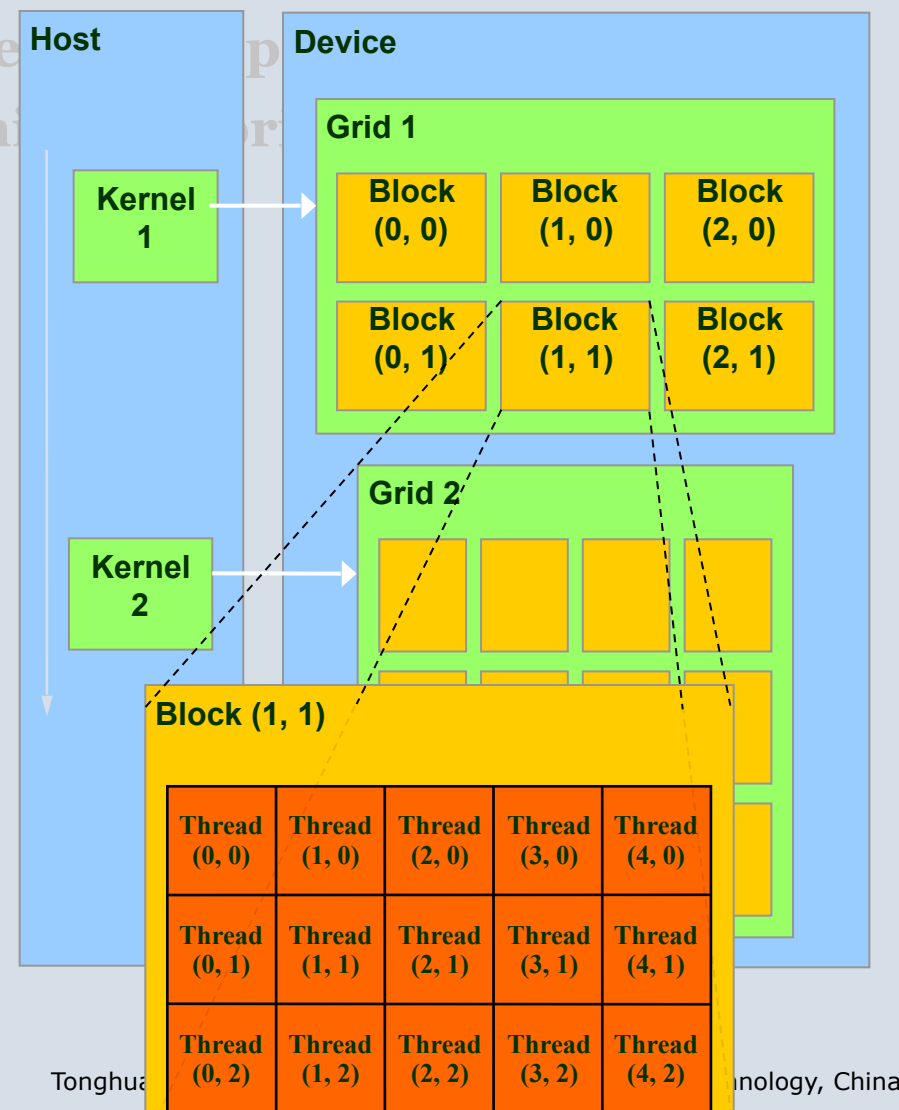
- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model for cooperating threads
- ✓ Lightweight synchronization primitives



# CUDA Abstraction

## •Key Parallel Abstractions in CUDA


- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model
- ✓ Lightweight synchronization




# Thread Hierarchy

## ● Thread → Block → Grid

```
void main(){  
    .....  
    int *dev_a,*dev_b,*dev_c;  
    .....  
    addKernel<<<1, 128>>>(dev_c, dev_a, dev_b);  
    .....  
}
```




```
void main(){  
    .....  
    int *dev_a,*dev_b,*dev_c;  
    .....  
    addKernel<<<100, 128>>>(dev_c, dev_a, dev_b);  
    .....  
}
```

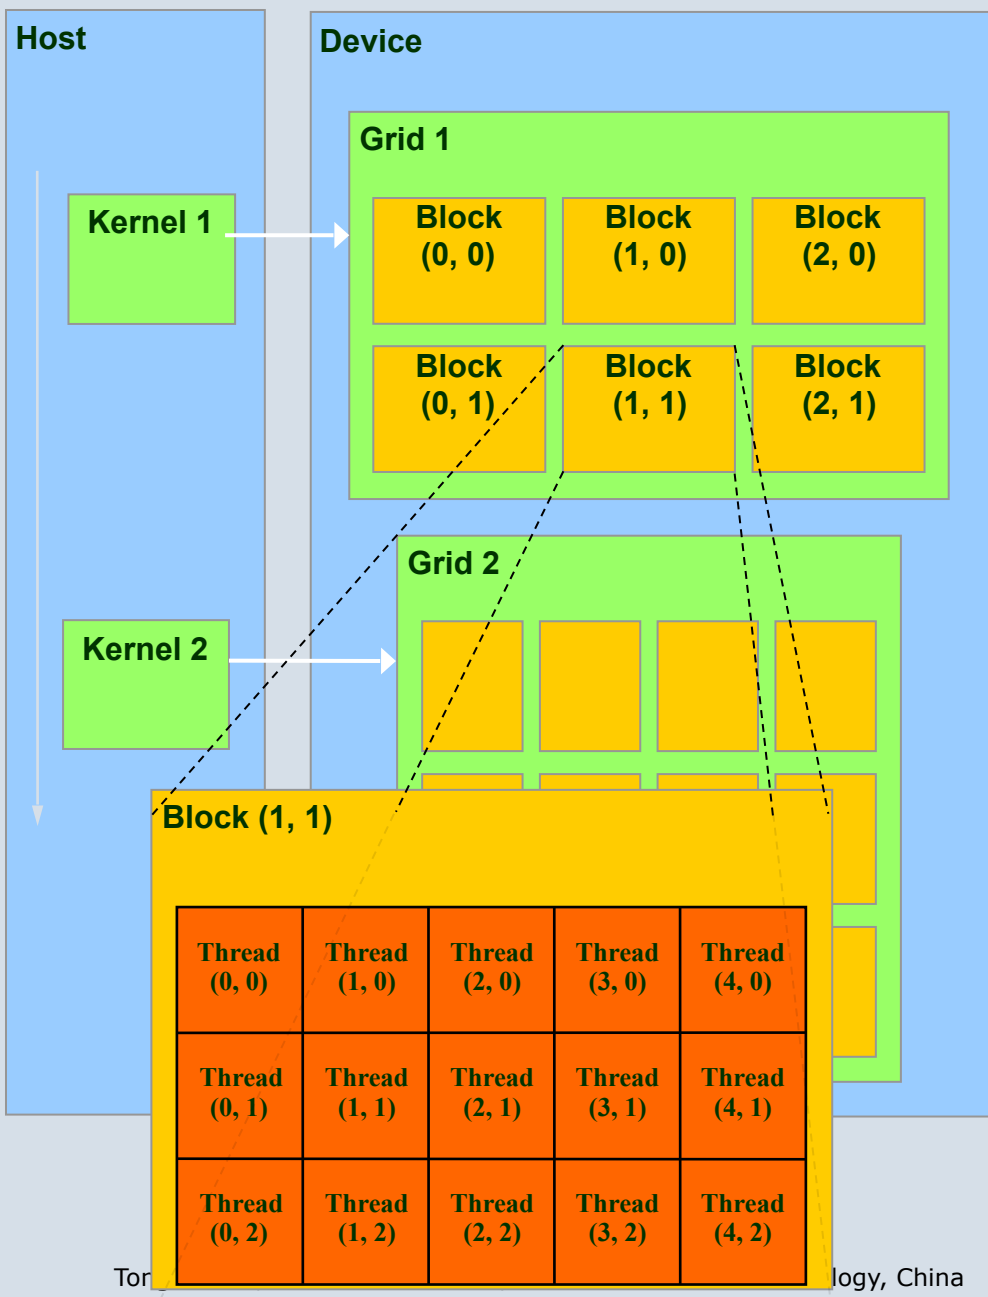


# Thread Hierarchy

## • Thread Mapping



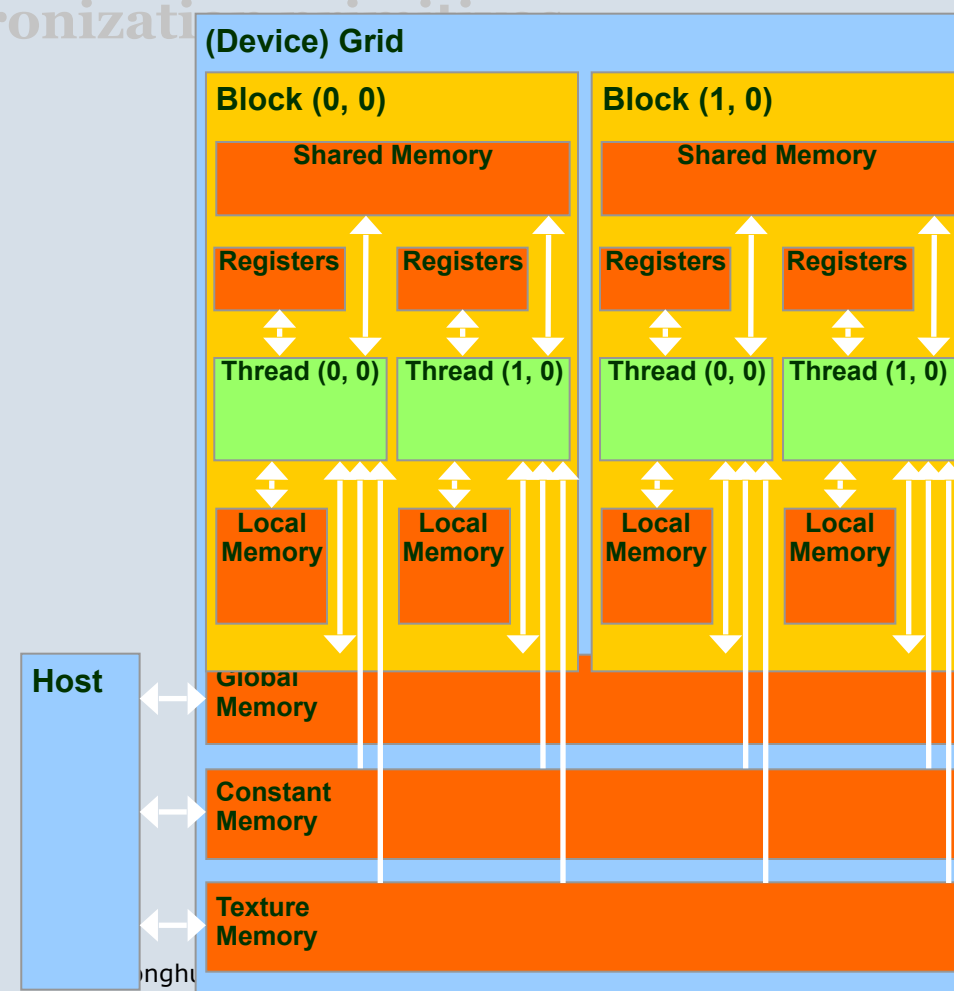
# Thread Hierarchy



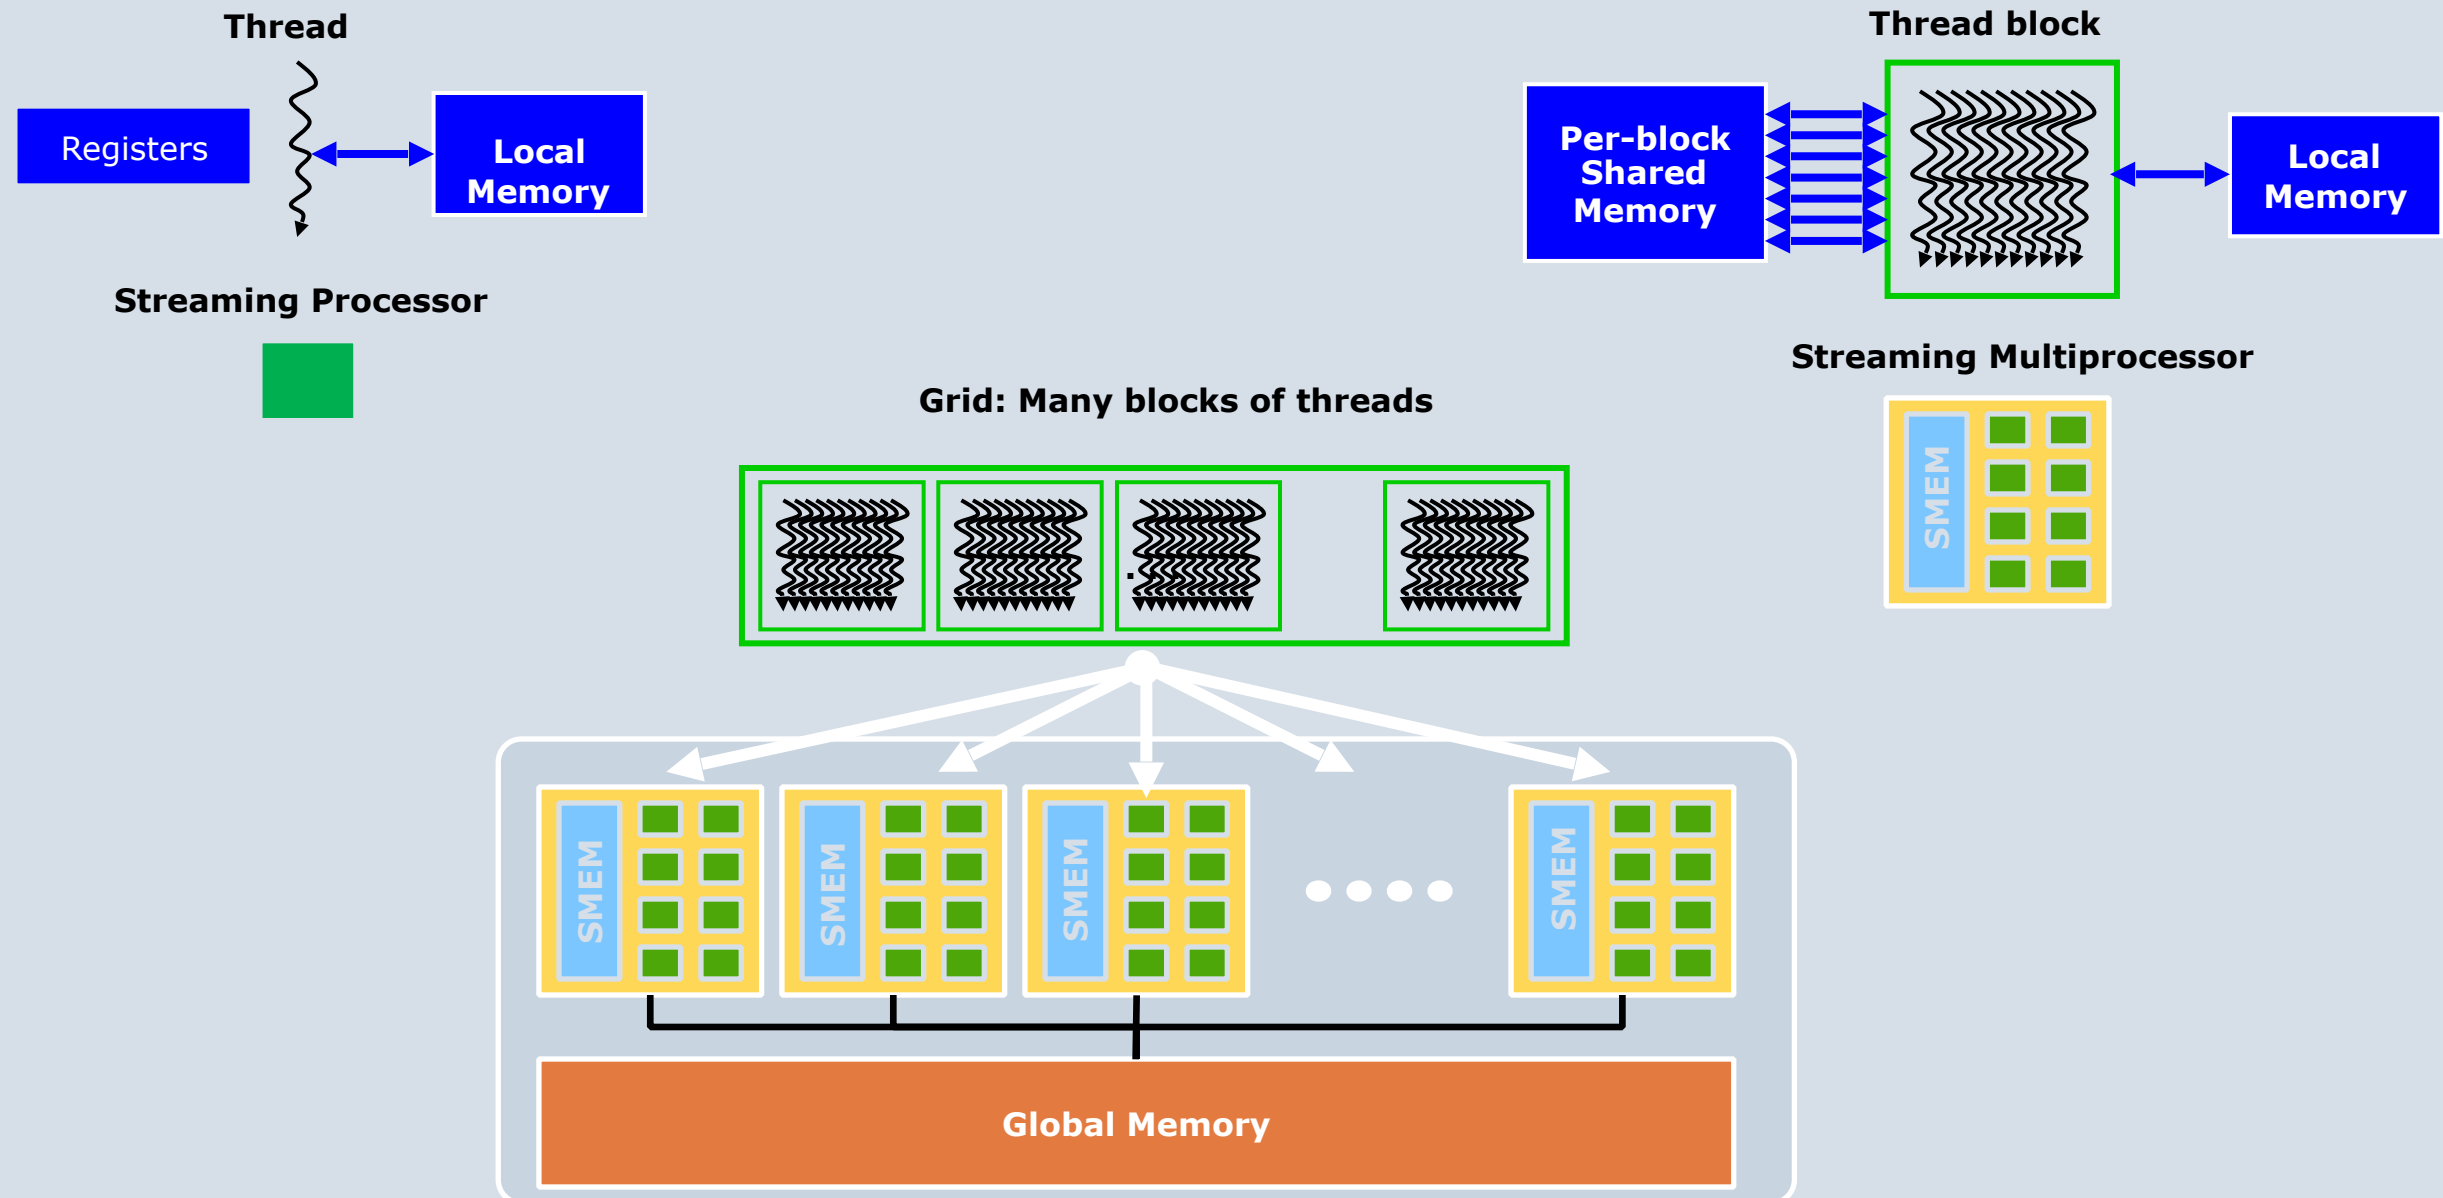
# CUDA Abstraction

## •Key Parallel Abstractions in CUDA

- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model for cooperating threads
- ✓ Lightweight synchronization primitives

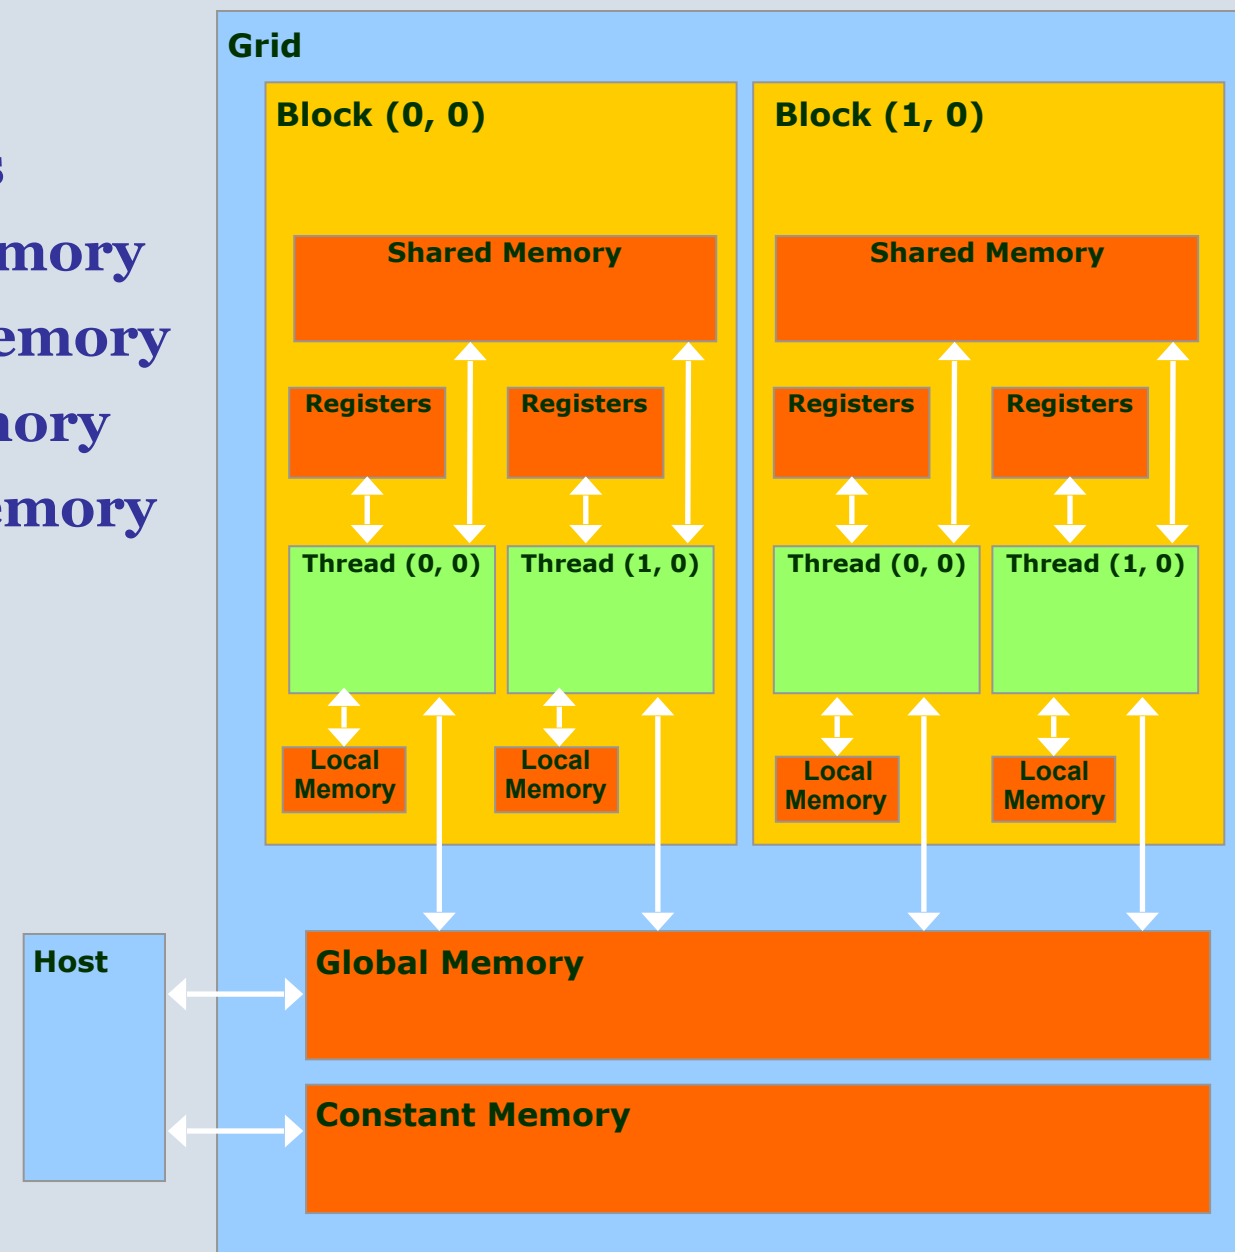


# Memory Model



# Memory Model

- Each thread can:
  - ✓ Read/write per-thread registers
  - ✓ Read/write per-thread local memory
  - ✓ Read/write per-block shared memory
  - ✓ Read/write per-grid global memory
  - ✓ Read/only per-grid constant memory





# GPU Multithreading Review

## ● Multithreaded GPU

- ✓ **CUDA core is multithreaded processor**
- ✓ **Run in group of 32 threads (called a warp)**
- ✓ **Zero-cost “context switching”**
  - Each thread has its own registers
- ✓ **Shared memory/L1 cache**

## ● Fine-grained multithreading

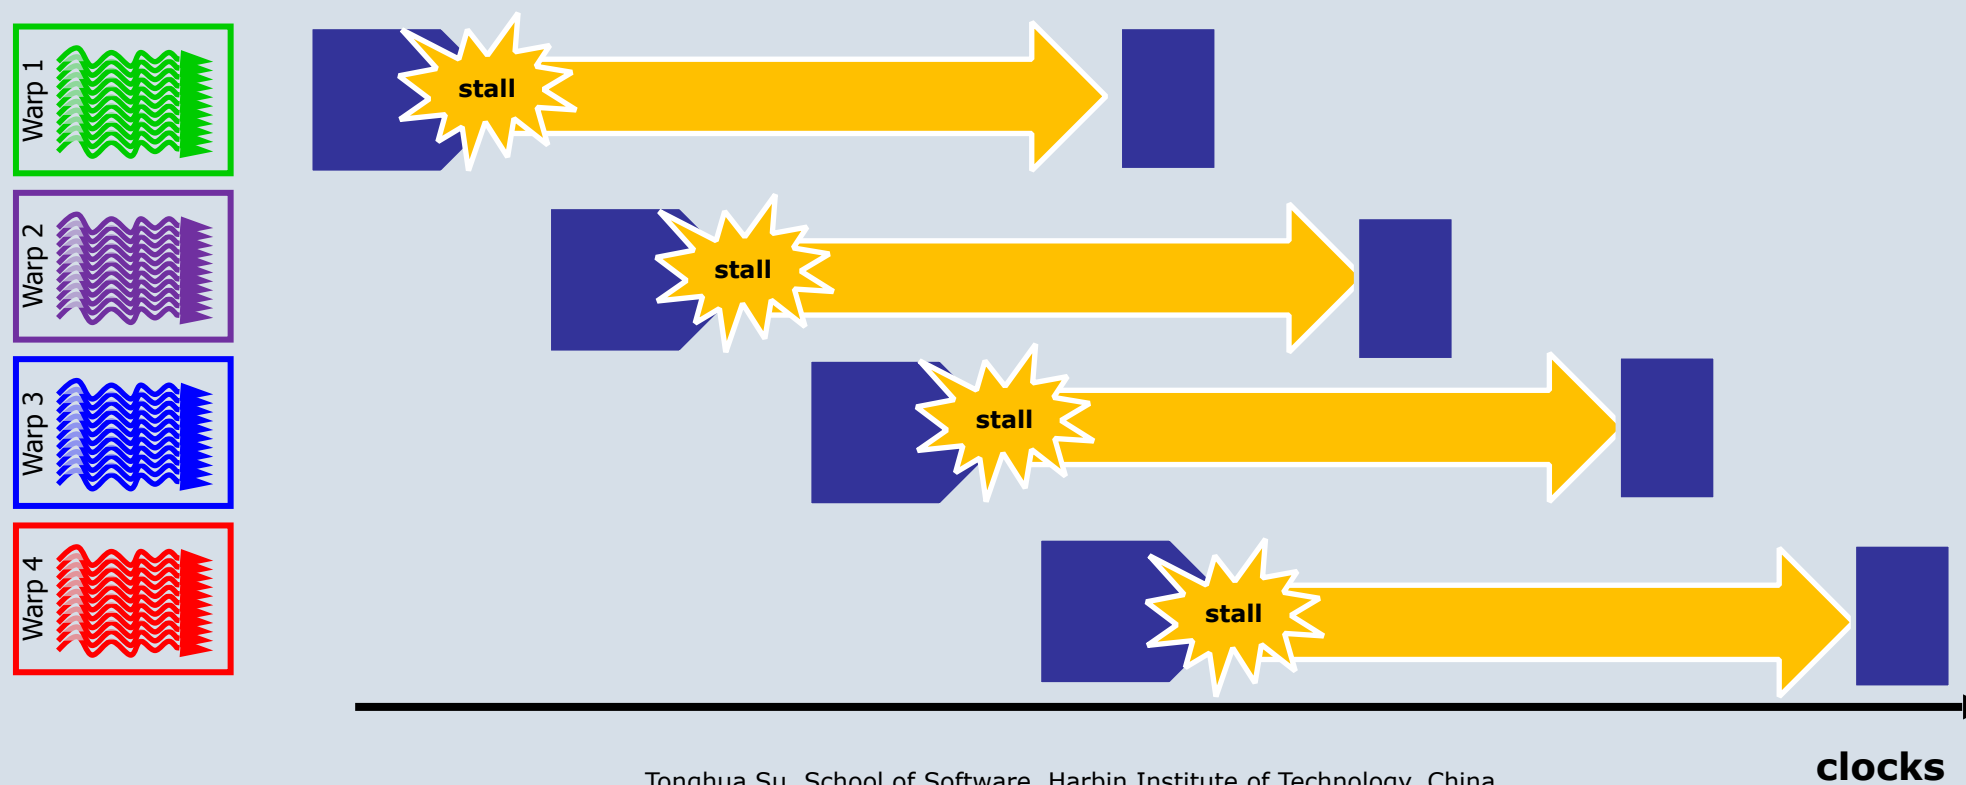
- ✓ **Able to switch between warps on each instruction**
- ✓ **Schedule without pre-emption**
- ✓ **Scheduling eligible warps in a round-robin fashion**



# Multithreading Review

## •Hiding Latency Stalls

- ✓ Fetch → Decode → Execute → **Memory** → Writeback
- ✓ Execution alternates between “active” warps, with warps becoming temporarily “inactive” when waiting for data
- ✓ **Lots of active warps** is the key to high performance




# CUDA Abstractions Review

- **CUDA is a parallel computing platform and programming model that enables dramatic increases in computing performance by harnessing the power of the GPU**
- **Key Parallel Abstractions in CUDA**
  - ✓ Hierarchy of concurrent threads
  - ✓ Shared memory model for cooperating threads
  - ✓ Lightweight synchronization primitives

# CUDA Abstractions Review

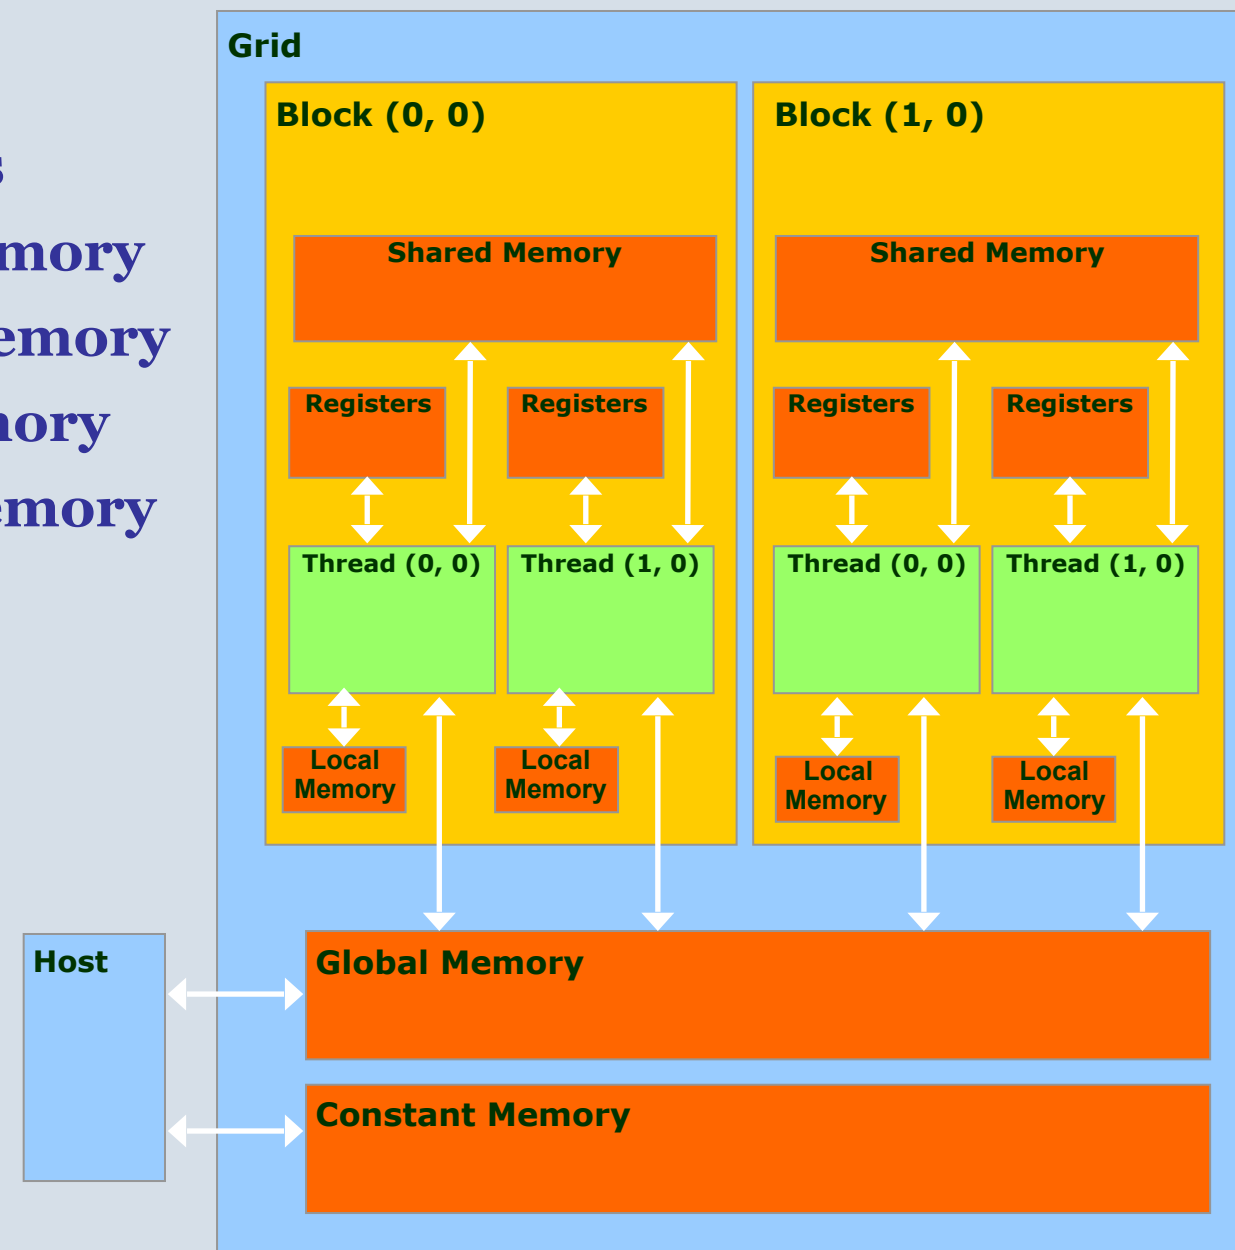
## • Thread Hierarchy and Thread Mapping



# CUDA Abstractions Review

## ● Memory Model

- ✓ Read/write per-thread **registers**
- ✓ Read/write per-thread **local memory**
- ✓ Read/write per-block **shared memory**
- ✓ Read/write per-grid **global memory**
- ✓ Read/only per-grid **constant memory**



# CUDA Abstraction

## ●Key Parallel Abstractions in CUDA

- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model for cooperating threads
- ✓ **Lightweight synchronization primitives**

# Synchronization

## ●Global Synchronization

- ✓ Finish a kernel and start a new one
- ✓ All writes from all threads complete before a kernel finishes

```
step1<<<grid1,blk1>>>(...);  
// The system ensures that all writes from step1  
complete.  
step2<<<grid2,blk2>>>(...);
```

- ✓ Would need to decompose kernels into before and after parts

# Synchronization

## •Threads Synchronization

✓ To ensure the threads visit the shared memory in order

✓ `__syncthreads()`

```
__device__ unsigned int count = 0;  
__shared__ bool isLastBlockDone;
```

```
__global__ void sum(const float* array, unsigned int N, volatile float* result)  
{  
    // Each block sums a subset of the input array.  
    float partialSum = calculatePartialSum(array, N);  
    if (threadIdx.x == 0) {  
        // Thread 0 of each block stores the partial sum to global memory.  
        result[blockIdx.x] = partialSum;  
        // Thread 0 makes sure the partial sum has been written to global memory.  
        __threadfence();  
        // Thread 0 signals that it is done.  
        unsigned int value = atomicInc(&count, gridDim.x);  
        // Thread 0 determines if its block is the last block to be done.  
        isLastBlockDone = (value == (gridDim.x - 1));  
    }  
    // make sure that each thread reads the correct value of isLastBlockDone.  
    __syncthreads();  
    if (isLastBlockDone) {  
        // The last block sums the partial sums stored in result[0 .. gridDim.x-1]  
        float totalSum = calculateTotalSum(result);  
        if (threadIdx.x == 0) {  
            result[0] = totalSum;  
            count = 0;  
        }  
    }  
}
```



# Synchronization

## ●Race Conditions

- ✓ What is the value of `a` in thread 0?
- ✓ What is the value of `a` in thread 127?

threadId:0	threadId:127
<code>// vector[0] was equal to 0</code>	
<code>vector[0] += 5;</code>	<code>vector[0] += 1;</code>
<code>...</code>	<code>...</code>
<code>a = vector[0];</code>	<code>a = vector[0];</code>

- ✓ CUDA provides **atomic** operations to deal with this problem

# Synchronization

## ●Atomics

- ✓ An atomic operation guarantees that only a single thread has access to a piece of memory while an operation completes
- ✓ Different types of atomic instructions:
  - `atomic{Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor}`
- ✓ Atomics are slower than normal load/store
- ✓ You can have the whole machine queuing on a single location in memory
- ✓ More types in Fermi
- ✓ Atomics unavailable on G80!
- ✓ e.g. `int atomicSub(int* address, int val);`

# Quiz

- 对输入的一张灰度图，要统计其灰度直方图，请编写CPU版代码
  - 可以通过OpenCV载入灰度图片，也可以用unsigned char数组gray模拟
  - 请采用线性内存存储数据
  - 暂且假定图片的长乘宽不大于1024

第1行数据



第2行数据



.....

最后1行数据



# Quiz

- 对输入的一张灰度图，要统计其灰度直方图，请编写可以完成此功能的**GPU**程序
  - 使用原子操作进行替换
  - 与**CPU**版代码结果对比

```
// Determine frequency of gray level in a picture
// gray have already been converted into unsigned char
// Each thread looks at one pixel and increments
// a counter atomically
__global__ void histogram(unsigned char* gray,
                          int* buckets)
{
    int i = threadIdx.x;
    unsigned char c = gray[i];
    buckets[c]++;
}
```